



CS 1550

Introduction and Lab 1

Jinpeng Zhou

jiz150@pitt.edu

Office Hours

- Office Hours (Zoom): Friday 11AM-3PM, and 7pm-9pm
- Email: jiz150@pitt.edu
- Piazza

Fall Term 2020-2021

Account
Dashboard
Courses
Calendar
Inbox
Commons
Help

Home
Announcements
Zoom
Panopto Video
Modules
Top Hat
Assignments
Piazza
Gradescope
People
NameCoach Roster
Grades

Upcoming Meetings Previous Meetings Personal Meeting Room Cloud Recording

☐ Show my course meetings only

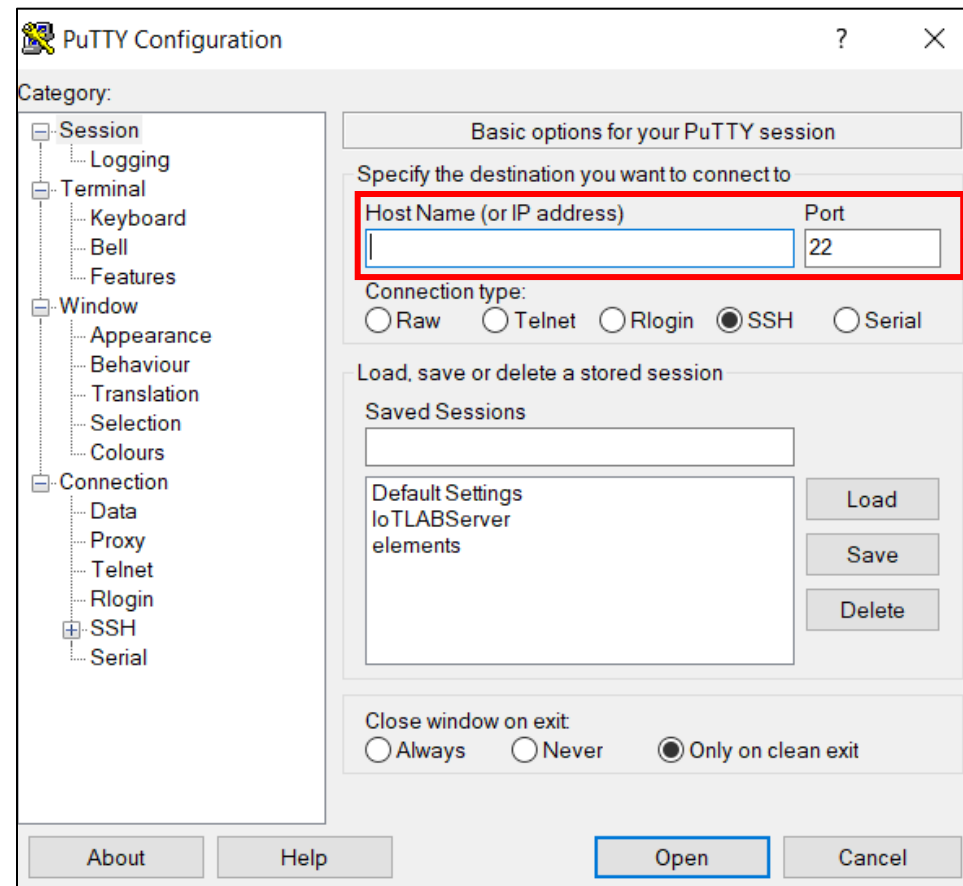
Start Time	Topic	Meeting ID
Tomorrow (Recurring) 10:00 AM	CS1550 Recitation 27114	961 5528 2833
Tomorrow (Recurring) 11:00 AM	CS1550 TA Office Hours	922 9846 3960
Tomorrow (Recurring) 3:30 PM	CS1550 Recitation 27113	915 5213 1465
Tomorrow (Recurring) 7:00 PM	CS1550 TA Office Hours	980 6716 7223

Use SSH to log into server

- Windows
 - **Putty**
- MacOS/Ubuntu
 - **Terminal**

Windows - Putty

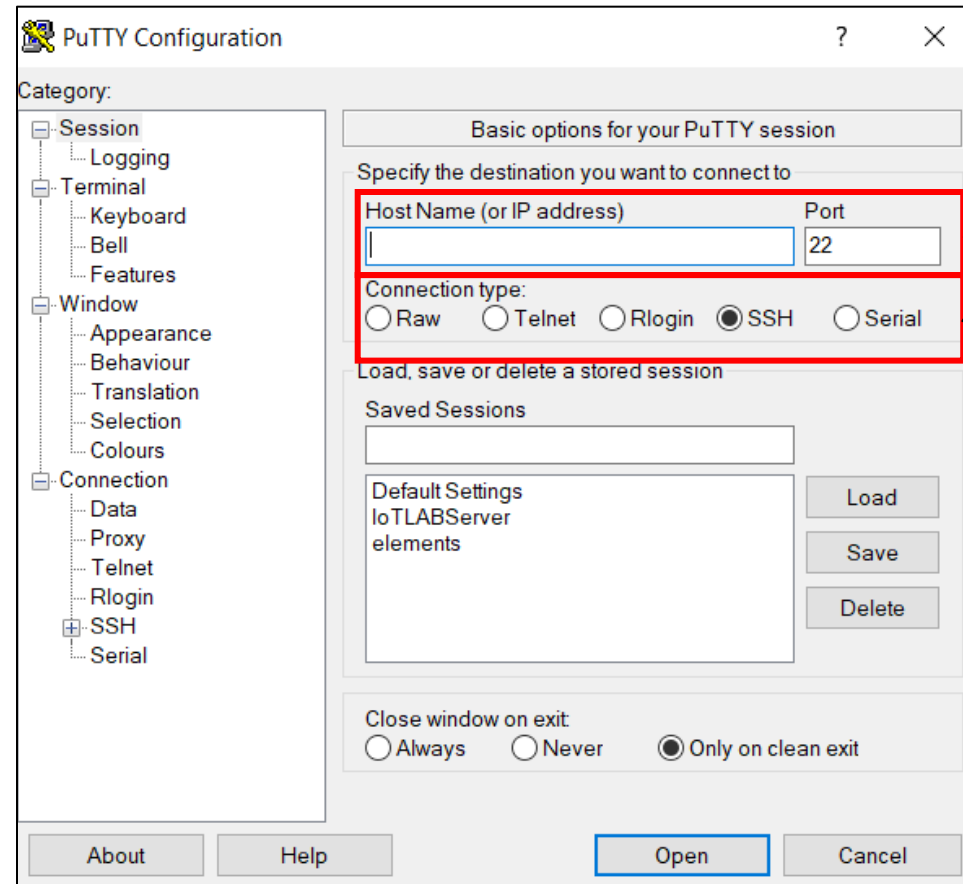
- Download from **www.putty.org**



server address:
linux.cs.pitt.edu at port 22

Windows - Putty

- Download from **www.putty.org**

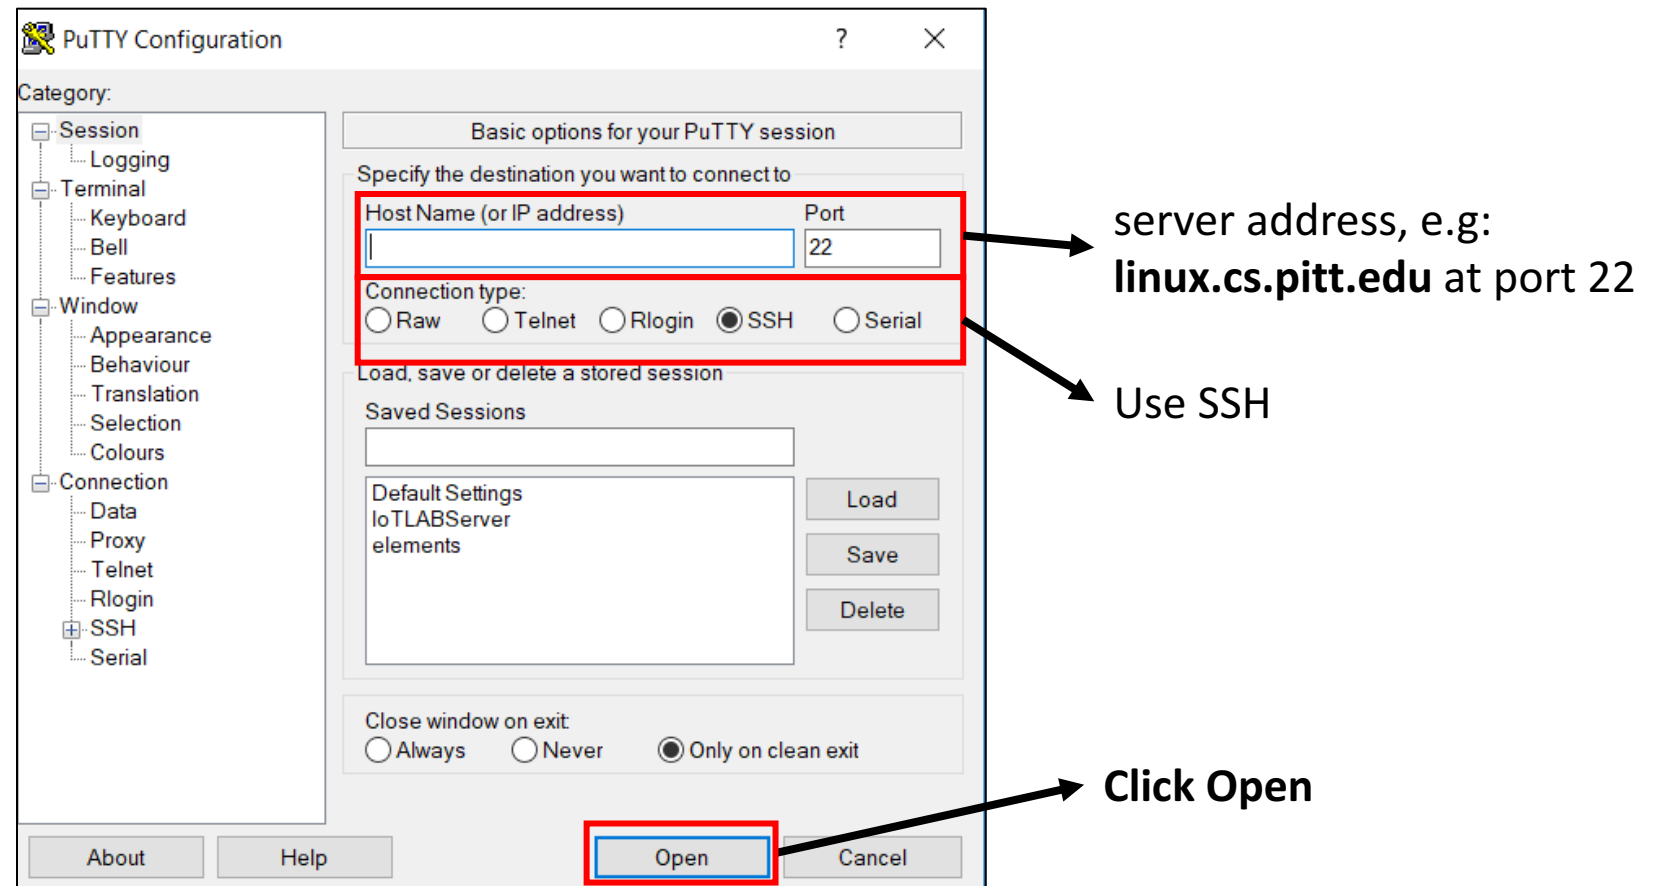


server address:
linux.cs.pitt.edu at port 22

Use SSH


Windows - Putty

- Download from **www.putty.org**



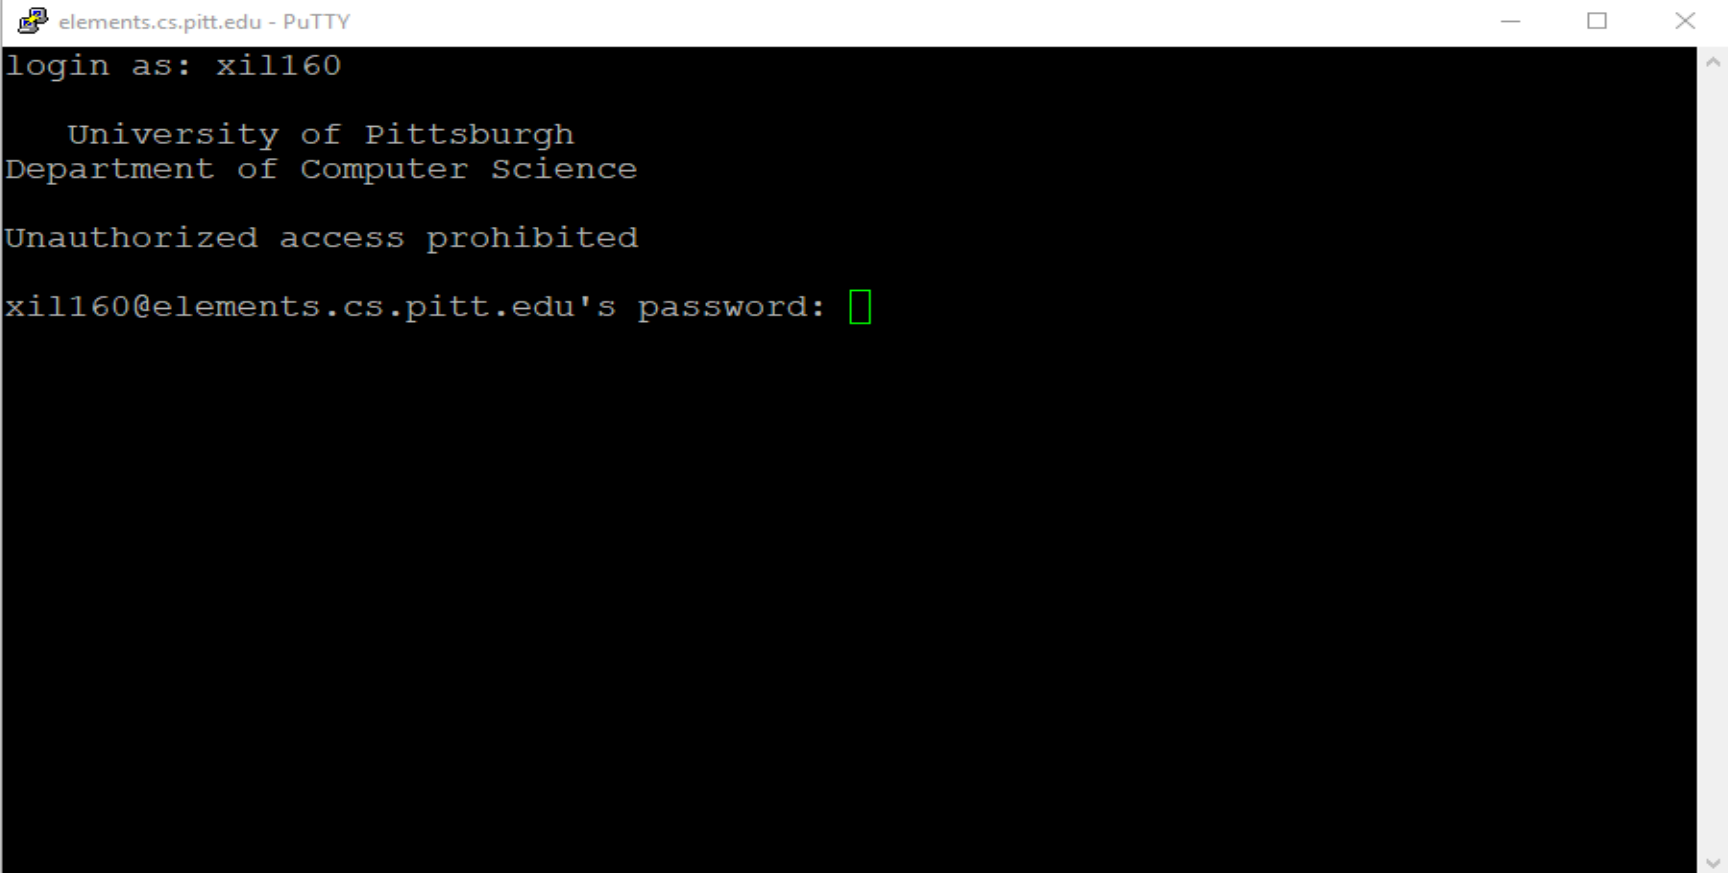
Windows - Putty

- Download from **www.putty.org**



Windows - Putty

- Download from **www.putty.org**



The screenshot shows a PuTTY terminal window titled "elements.cs.pitt.edu - PuTTY". The terminal output is as follows:

```
login as: xill60

  University of Pittsburgh
Department of Computer Science

Unauthorized access prohibited

xill60@elements.cs.pitt.edu's password: [redacted]
```

The password field is represented by a series of green squares, indicating that the password has been entered but is not visible.

MacOS/Ubuntu - Terminal

ssh command

User name

```
[zjp:~$ ssh jiz150@linux.cs.pitt.edu
```

```
University of Pittsburgh  
Department of Computer Science  
  
Unauthorized access prohibited
```

```
jiz150@linux.cs.pitt.edu's password:
```



password

THIS SYSTEM IS FOR THE USE OF AUTHORIZED USERS ONLY.

Individuals using this computer system without authority, or in excess of their authority, are subject to having all of their activities on this system monitored and recorded by system personnel.

In the course of monitoring individuals improperly using this system, or in the course of system maintenance, the activities of authorized users may also be monitored.

Anyone using this system expressly consents to such monitoring and is advised that if such monitoring reveals possible evidence of criminal activity, system personnel may provide the evidence of such monitoring to law enforcement officials.

(1) thompson \$

Basic Linux Shell commands

- Check Current Directory – **pwd**
- List directories - **ls**
- Create/Remove directory – **mkdir/rmdir**
- Remove files – **rm**
- Copy files from anywhere to anywhere – **cp**
 - **cp <current path> <new path>**
 - **cp some_text.txt Desktop/**
- Move files from anywhere to anywhere – **mv**
 - **mv <current path> <new path>**
 - **mv some_text.txt Desktop/**

Basic Linux Shell commands

- Manual page

- Type “man COMMAND” to see the manual page for a command

- E.g., type “man rm” to see how to use “rm”. Type “q” to exit the man page

- Possible problem:

- IF: you saw “No manual entry” on linux.cs.pitt.edu

```
(1) thompson $ man rm  
No manual entry for rm
```

- THEN: edit the “~/.bash_profile” file by adding the “export” line in the end of the file, then run “source ~/.bash_profile”

```
(2) thompson $ vim ~/.bash_profile
```

```
export MANPATH="$MANPATH:/usr/share/man"
```

```
(3) thompson $ source ~/.bash_profile
```

- Or you can simply search something like “man page rm” online

Building projects

- Small programs
 - single file

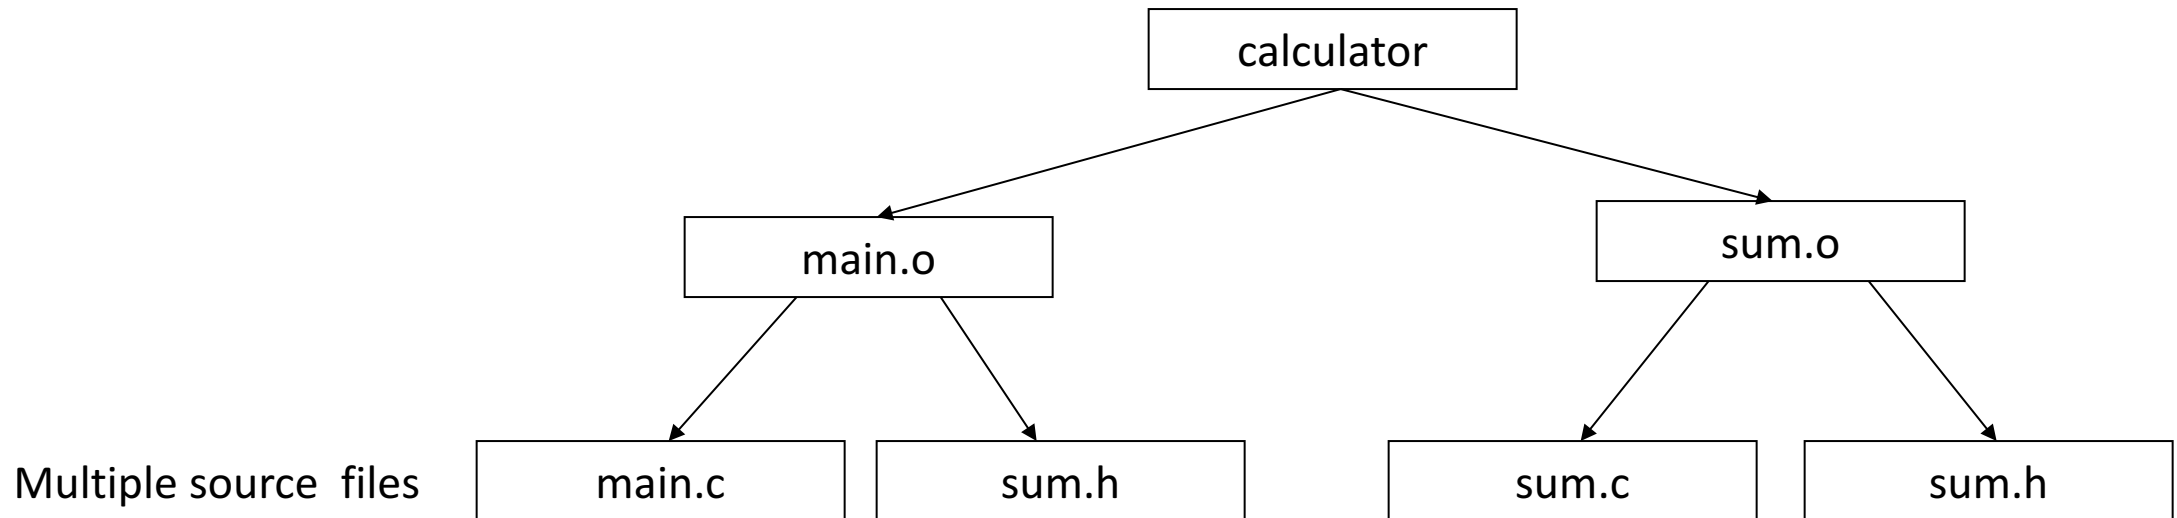
Building projects

- Small programs (easy to compile)
 - single file

```
~home-laptop->$ gcc main.c -o calculator
```

Building projects

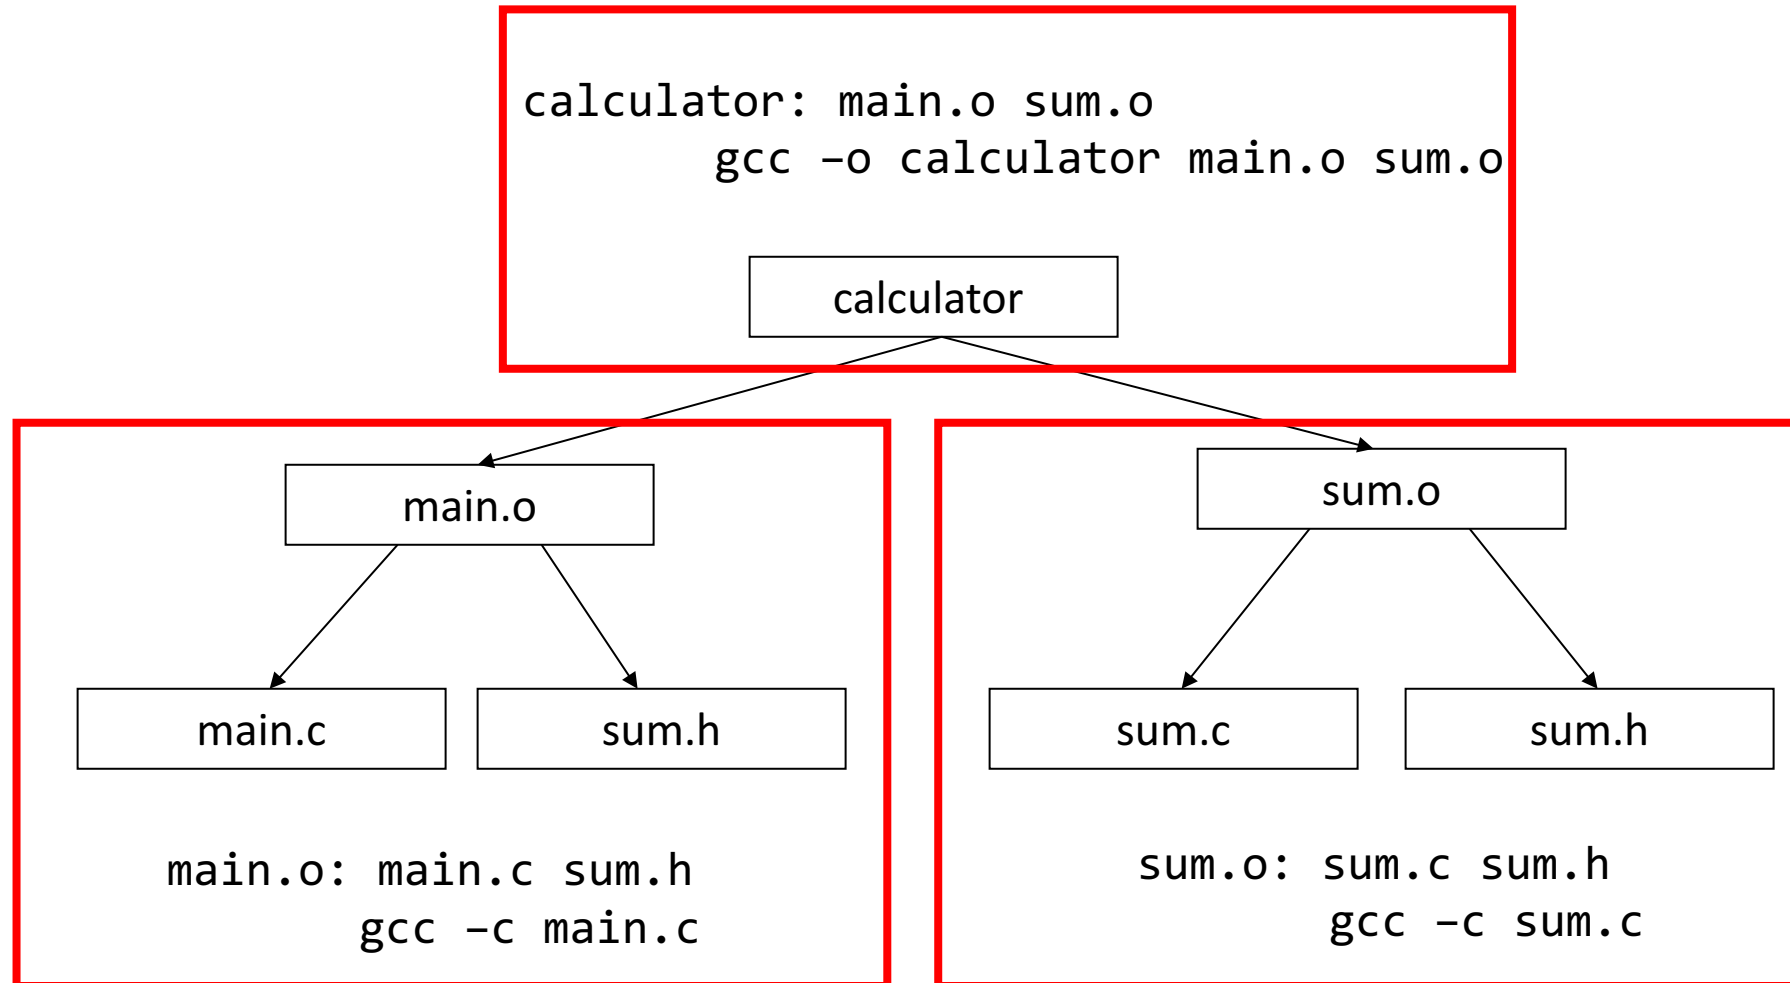
- Small programs (easy to compile)
 - single file
- Bigger programs
 - multiple files



Use the Makefile

TARGET: DEPENDENCIES
ACTION

Use the Makefile



Makefile example

calculator: `main.o` `sum.o`

`gcc -o calculator main.o sum.o`

`main.o:` `main.c` `sum.h`

`gcc -c main.c`

`sum.o:` `sum.c` `sum.h`

`gcc -c sum.c`

The Makefile

- Running “make”

```
~home-laptop->$ make calculator
```

The Makefile

```
calculator: main.o sum.o
    gcc -o calculator main.o sum.o
main.o: main.c sum.h
    gcc -c main.c
sum.o: sum.c sum.h
    gcc -c sum.c
clean:
    rm -f *.o
```

The Makefile

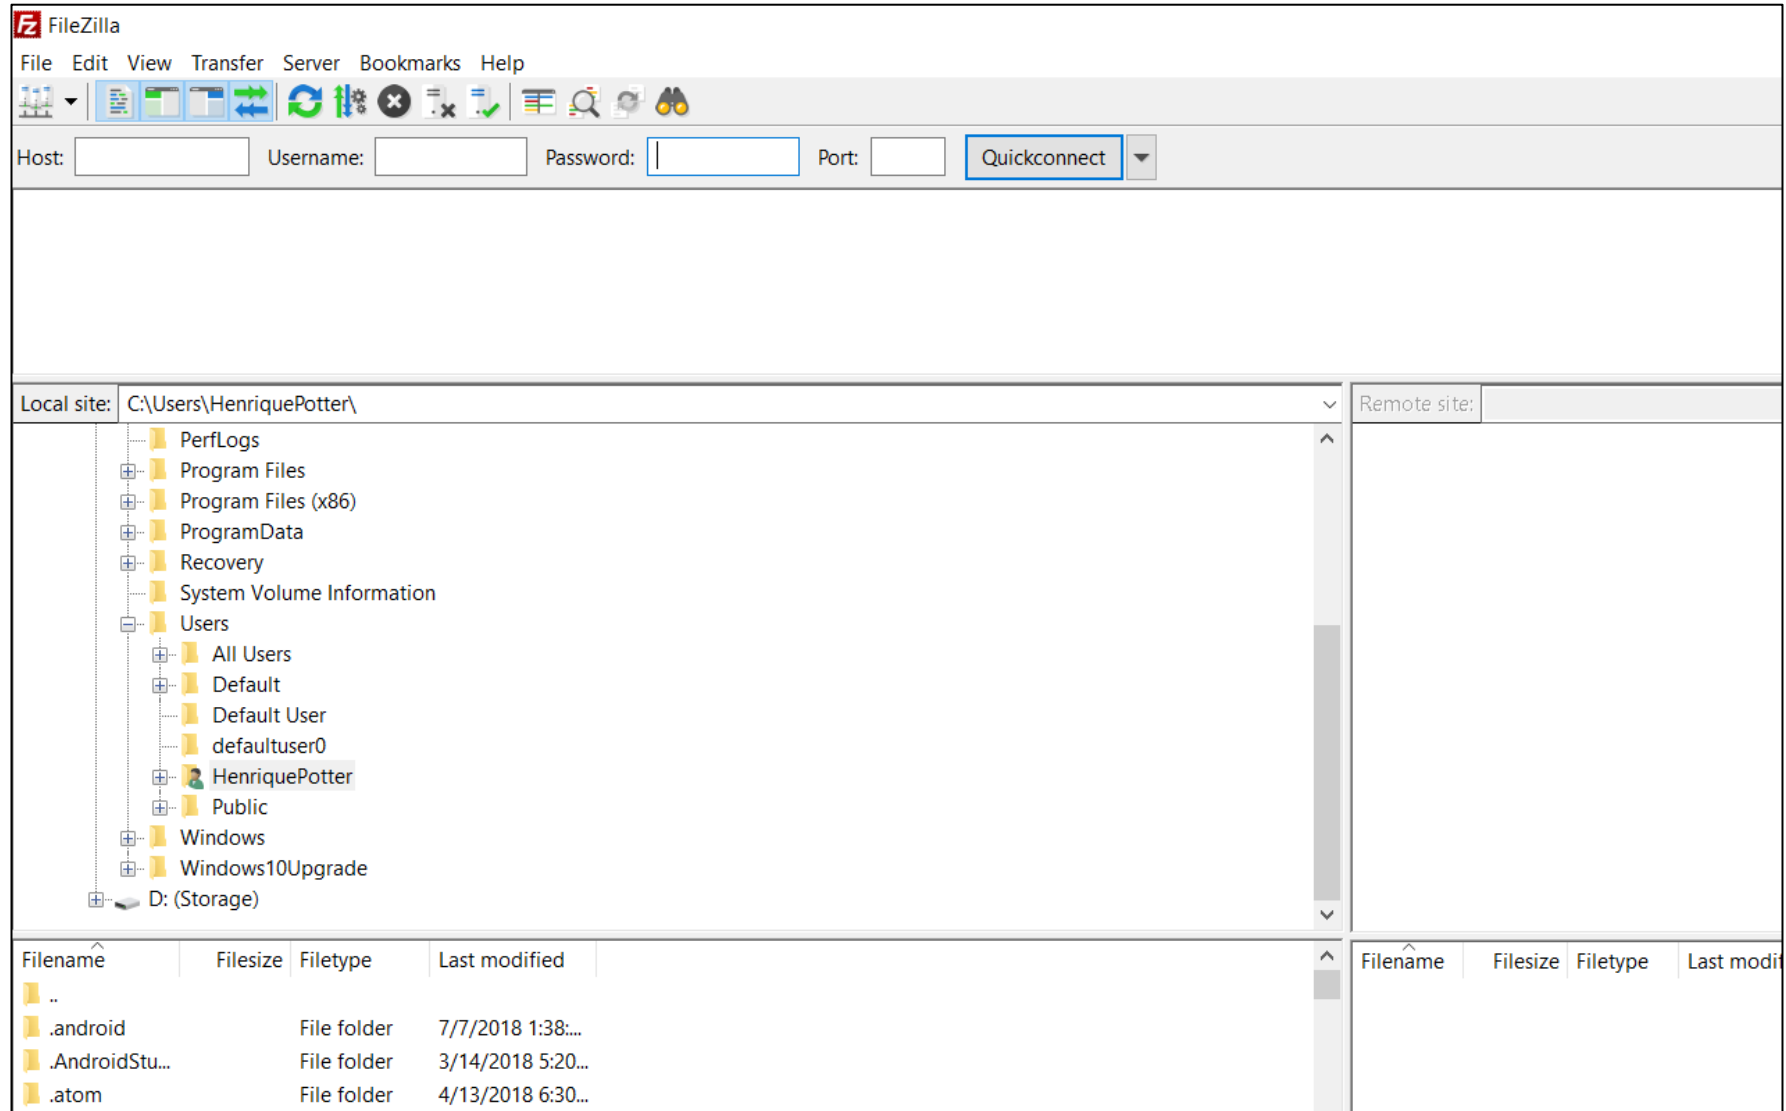
- Running “make”

```
~home-laptop->$ make clean
```

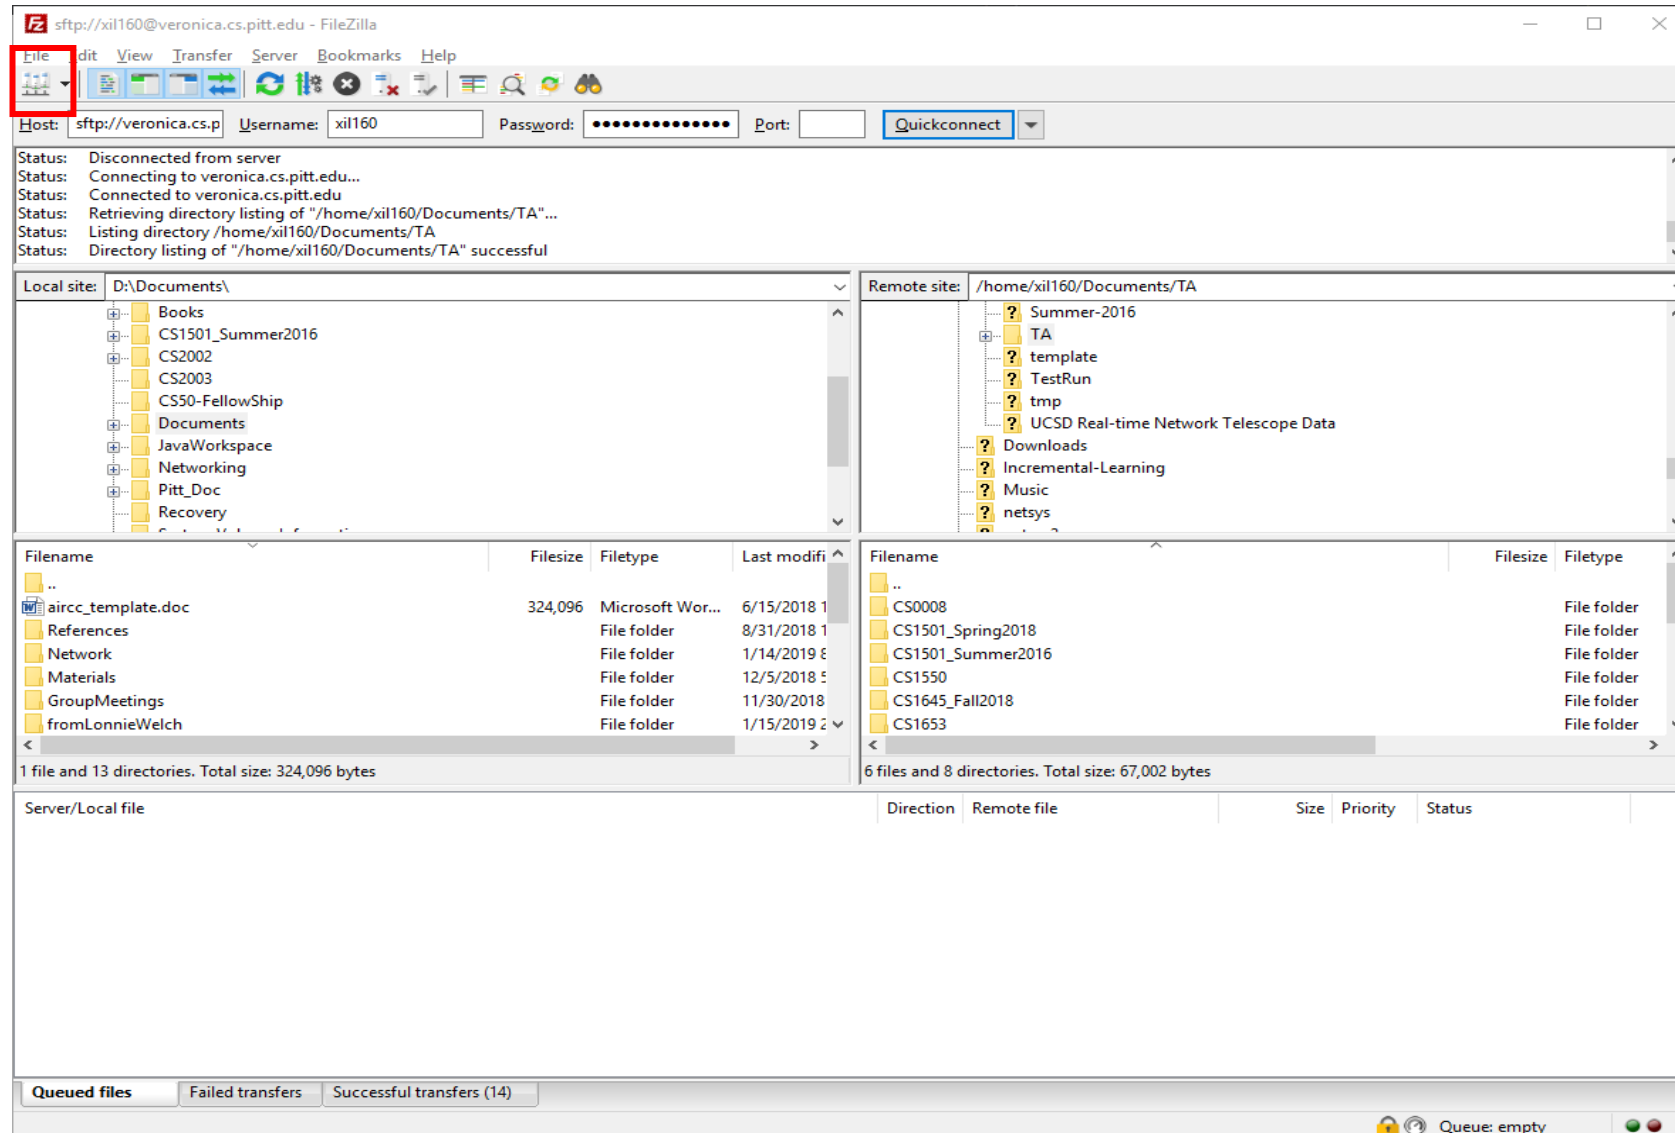
Upload to / download from server

- FileZilla – Windows/MacOS
 - GUI based
- Or use the “scp” command
 - Try check scp’s man page

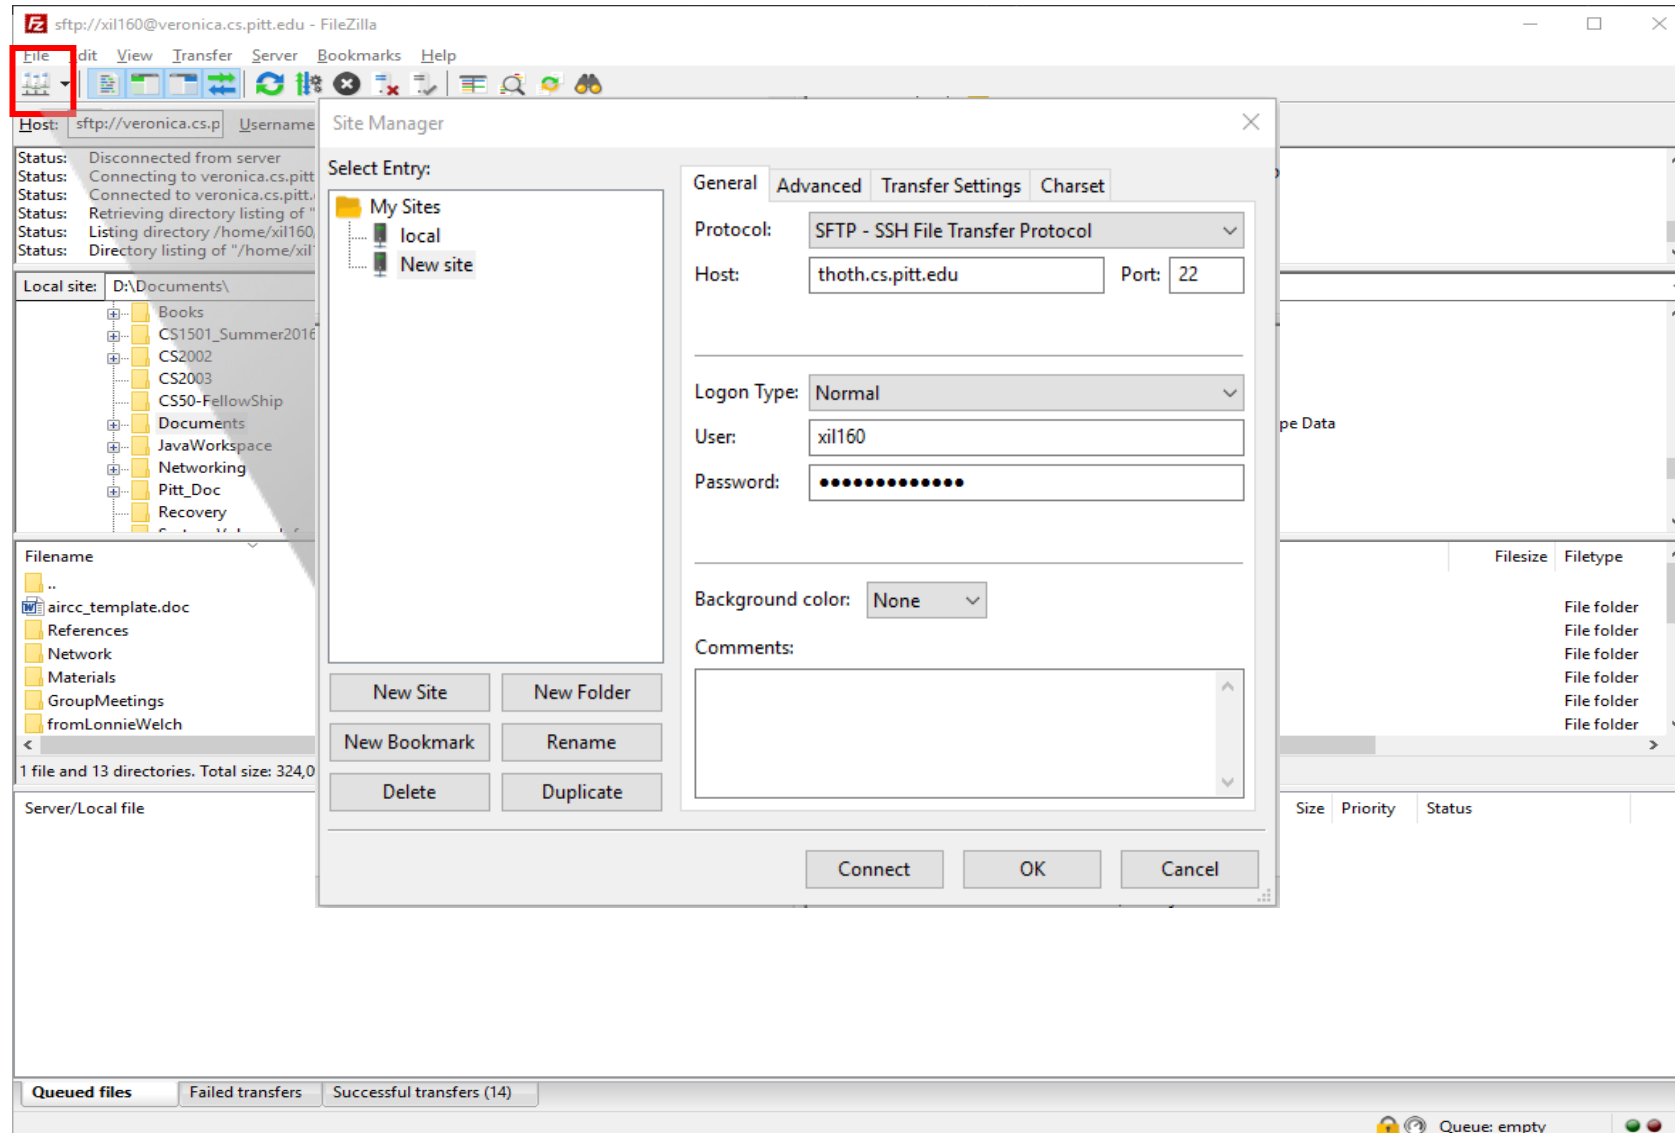
GUI based FTP Clients



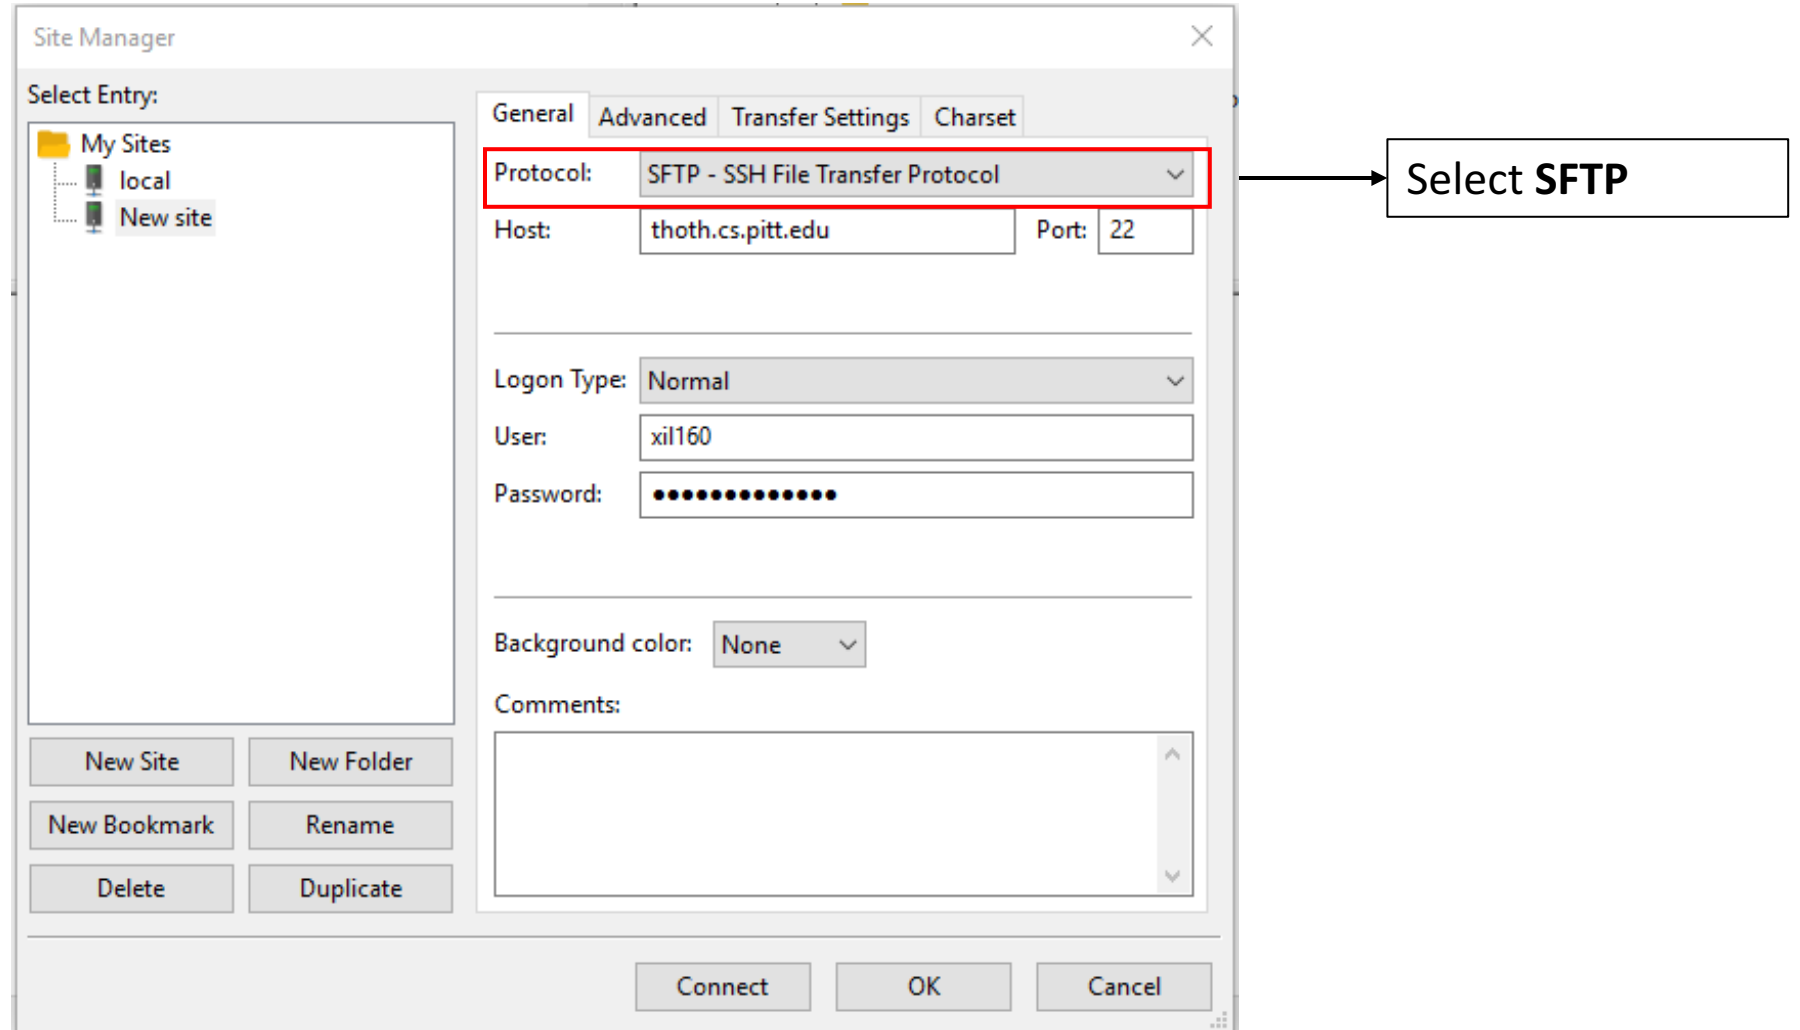
GUI based FTP Clients



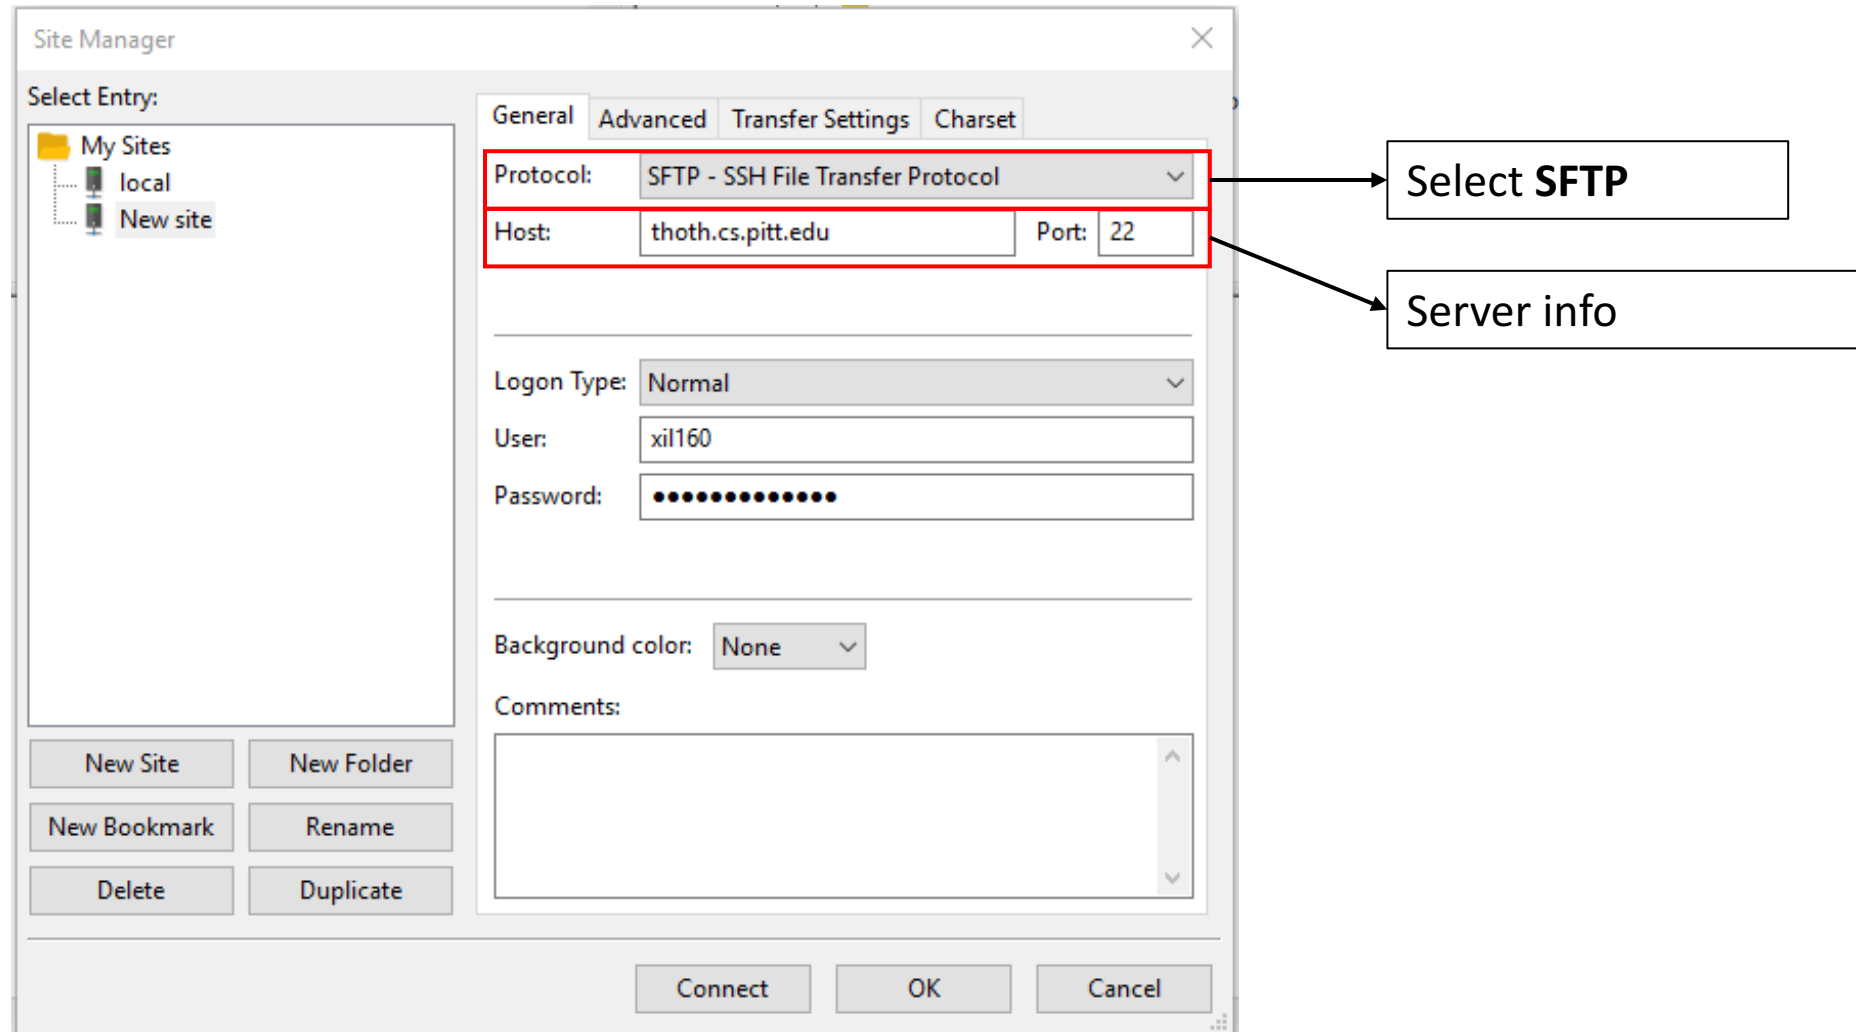
GUI based FTP Clients



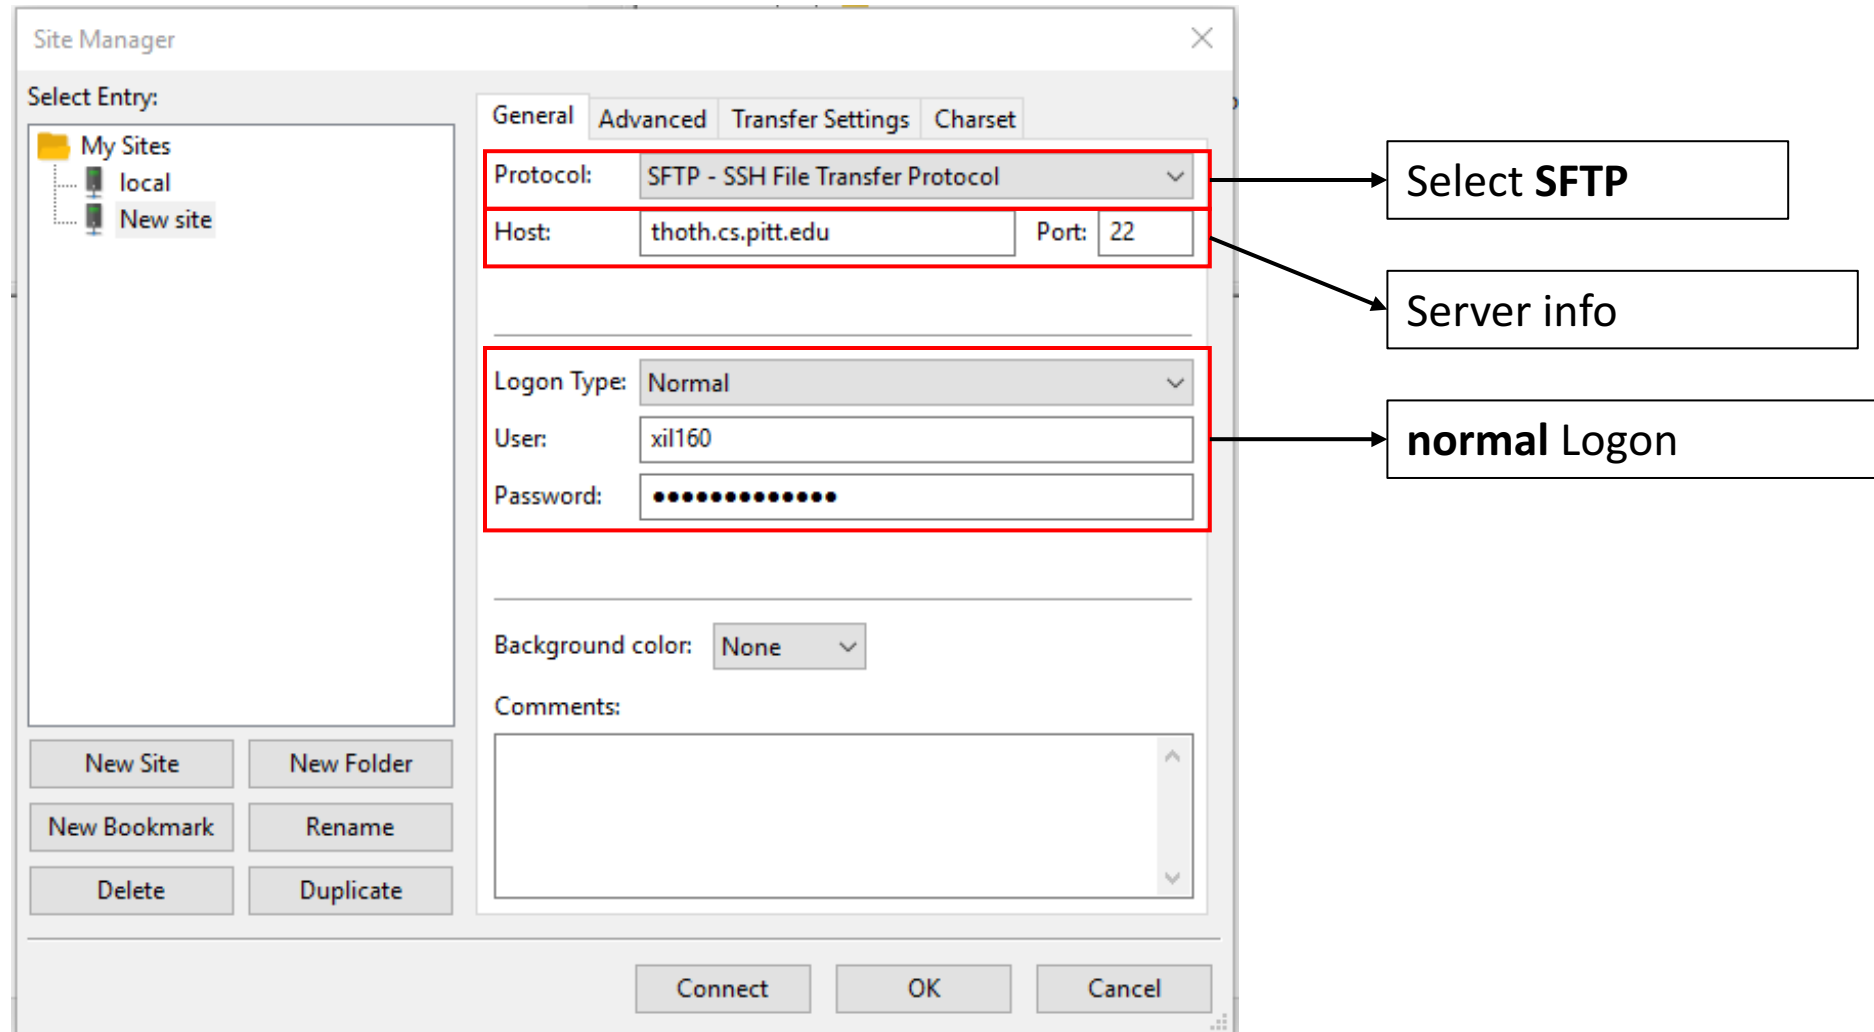
GUI based FTP Clients



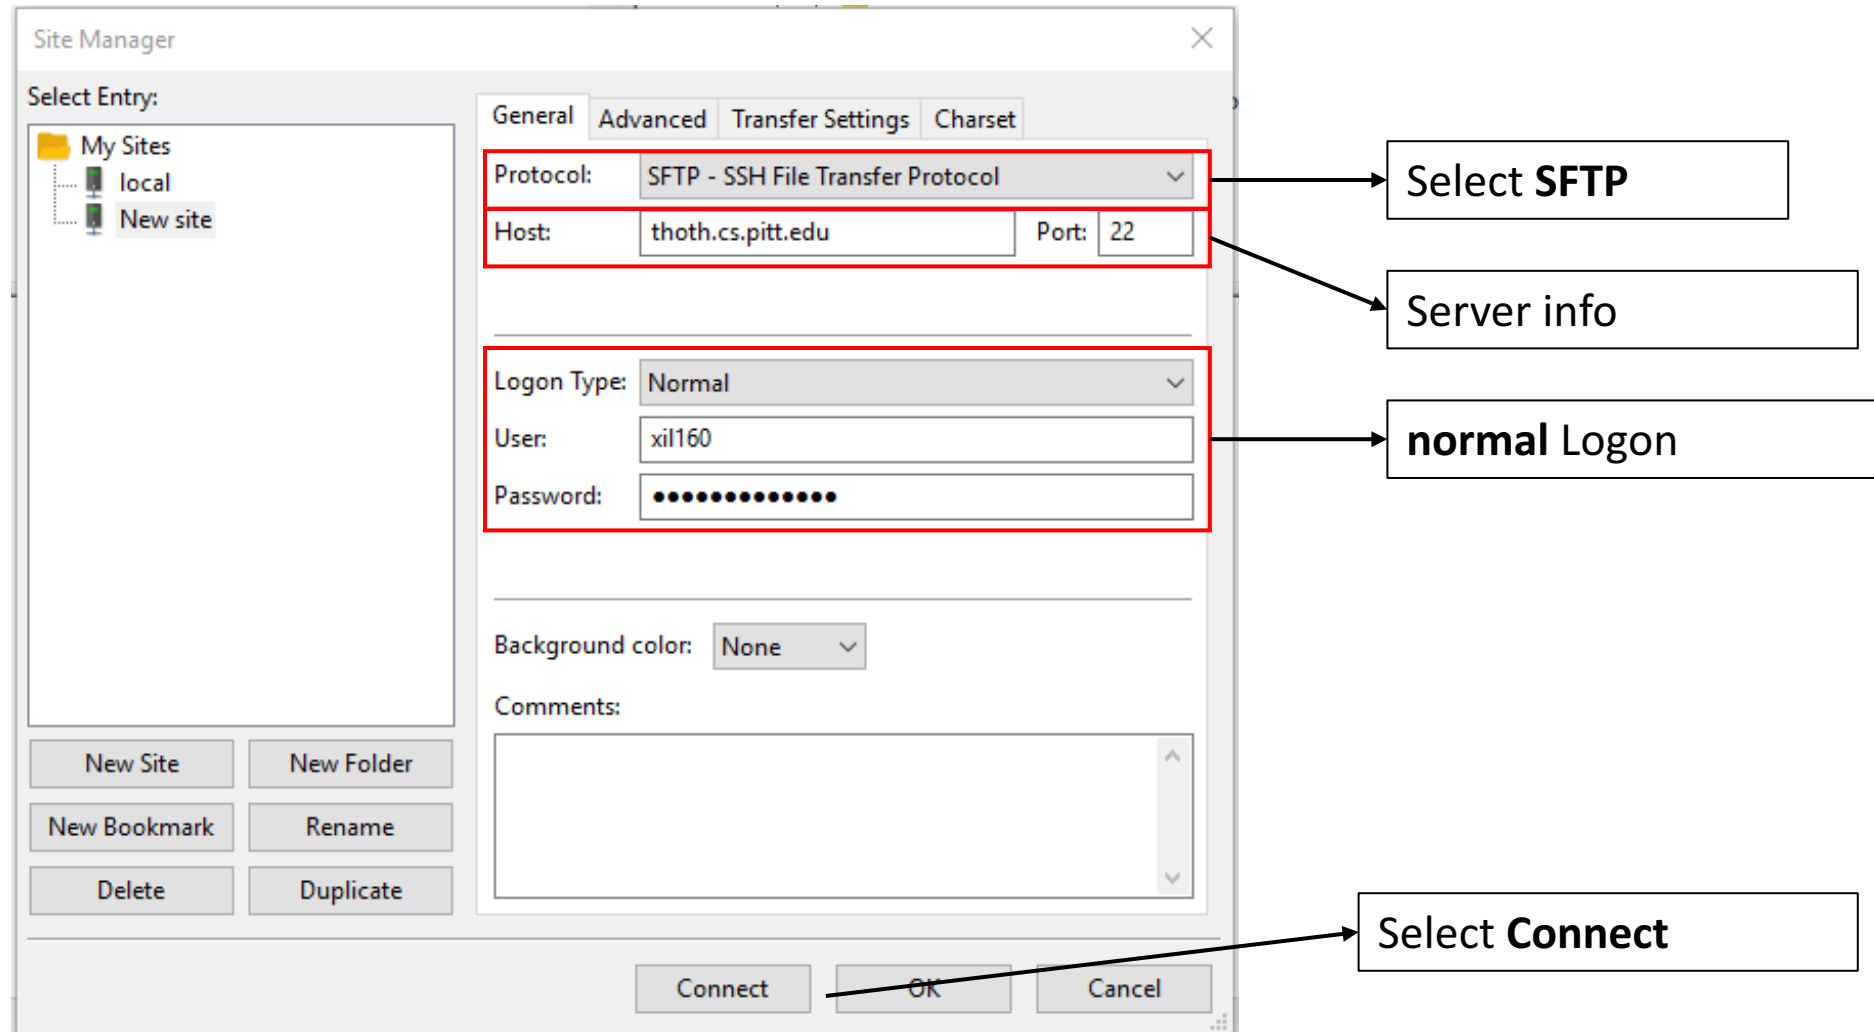
GUI based FTP Clients



GUI based FTP Clients




GUI based FTP Clients



Using Atom

- Good coding tool
 - Easily search for variables names in any file inside a folder
 - Lots of plugins
- Synchronize Files
 - Install ***remote-sync*** (recommended)
 - Right click main project folder
 - Navigate to Remote Sync > Configure
 - Fill in the details / select options
 - Hit save
 - <https://atom.io/packages/remote-sync>



C Pointers

- Pointers are variables that contain *memory addresses* as their values.
- A variable name *directly* reference a value.
- A pointer *indirectly* reference a value.
- A pointer variable must be declared before it can be used.

C Pointers

- Examples of pointer declarations:
 - `int *a;` * tells the compiler that the variable is to be a pointer, and the type of data that the pointer points to
 - `float *b;`
- & and *:
 - & -- “address operator” gives or produces the memory address of a data variable.
 - * -- “dereferencing operator” provides the contents in the memory location specified by a pointer.

C Pointers

- A simple Example

```
int number;
```

```
int *ptr1;
```

```
int *ptr2;
```

```
ptr1 = &number;
```

```
number = 2;
```

```
ptr2 = &number;
```

```
printf("\n*ptr1 = %d\t*ptr2 = %d\n", *ptr1, *ptr2);
```


C Pointers

- Arrays

```
int demoArray[5] = {8, 19, 34, 0, 3};
```

```
printf("Element \t Address \t Value \n");
```

```
for (int i = 0; i < 5; i++) {
```

```
    printf("demoArray[%d]\t%p\t%d\n", i, &demoArray[i], demoArray[i]);
```

```
}
```

```
//array names are just pointers to the first Element
```

```
printf("\ndemoArray\t%p\n", demoArray);
```

```
//dereference
```

```
printf("\n*demoArray\t%d\n", *demoArray);
```

```
printf("\n*(demoArray+2)\t%d\n", *(demoArray+2));
```



System Call - exercise

CS 1550 – xv6

- Simple Unix-like **operating system** for teaching, developed in 2006.
 - Provides basic services for running programs
 - https://pdos.csail.mit.edu/6.828/2012/xv6.html?utm_source=dlvr.it&utm_medium=tumblr

A black rectangular box containing the text "xv6 OS" in a yellow, monospace-style font.

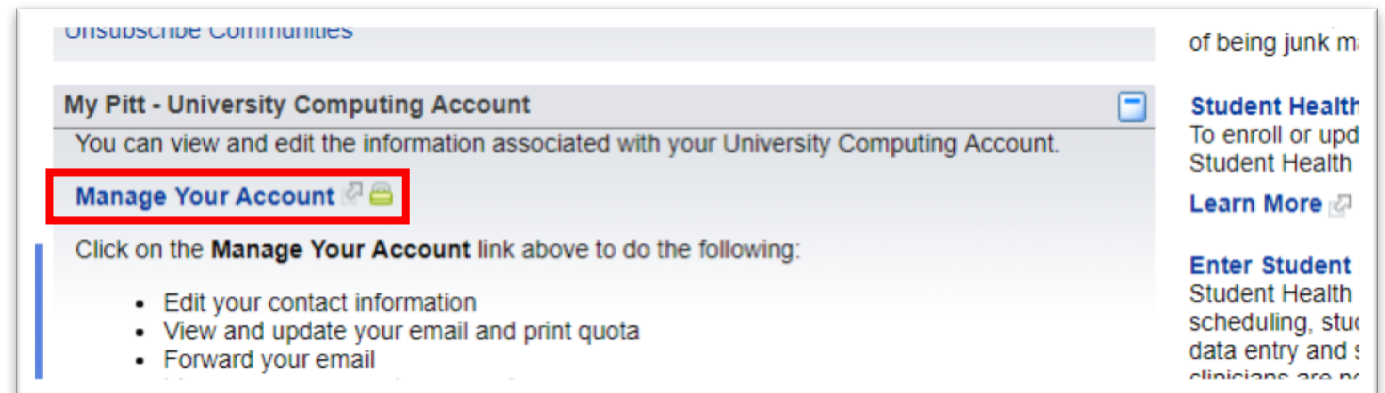

xv6 OS

CS 1550 – xv6

- Has a **subset of traditional** system calls
 - **fork()** Create process
 - **exit()** Terminate current process
 - **wait()** Wait for a child process
 - **kill(pid)** Terminate process pid
 - **getpid()** Return current process's id
 - **sleep(n)** Sleep for n time units
 - **exec(filename, *argv)** Load a file and execute it sbrk(n)
 -

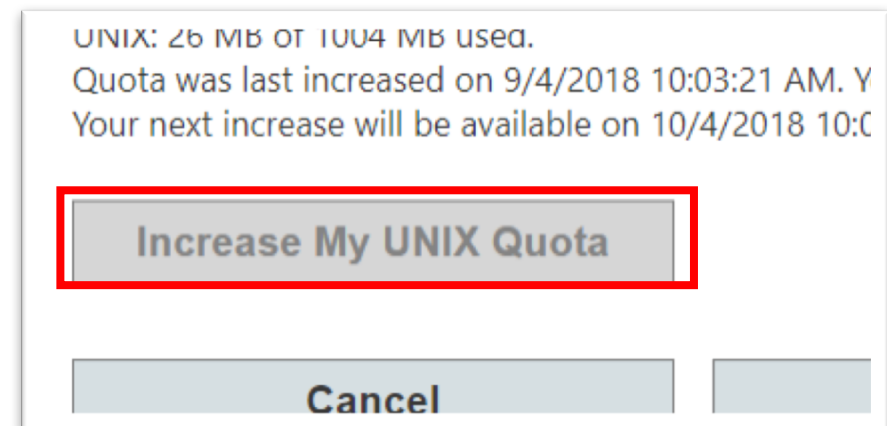

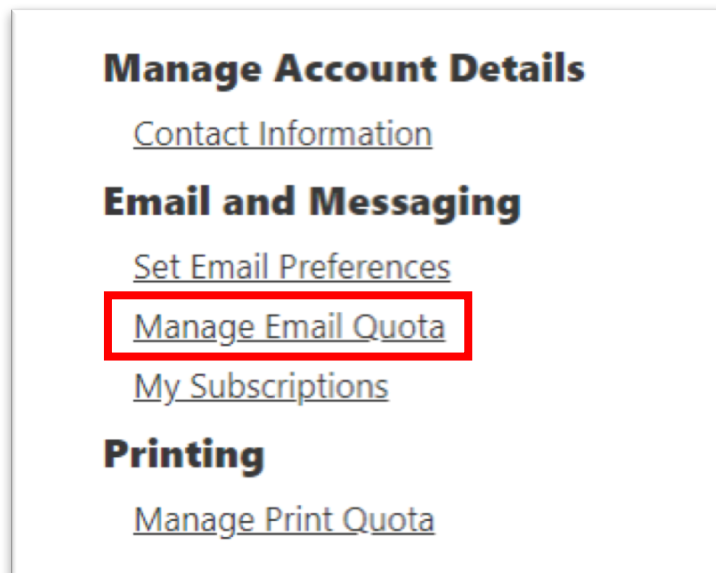
CS 1550 – Compile and Run xv6

1. Extend disk Quota, if you have less then 500mb free space
 - a) Log in to <https://my.pitt.edu>
 - b) Click on "Profile" at the top-right corner of the screen
 - c) Click on "Manage Your Account"
 - d) Click on "EMAIL & MESSAGING" -> "UNIX QUOTA"
 - e) Click on "Increase My UNIX Quota"




CS 1550 – Compile and Run xv6

1. Extend disk Quota, if you have less then 500mb free space
 - a) Log in to <https://my.pitt.edu>
 - b) Click on "Profile" at the top-right corner of the screen
 - c) Click on "Manage Your Account"
 - d) Click on "EMAIL & MESSAGING" -> "UNIX QUOTA"
 - e) Click on "Increase My UNIX Quota"



CS 1550 – xv6

- Log in to **linux.cs.pitt.edu**
 - **ssh user_name@linux.cs.pitt.edu**
- Download the xv6 source code from github
 - **git clone git://github.com/mit-pdos/xv6-public.git**
- Get into the cloned xv6 source code folder
 - **cd xv6-public**
- Compile and run the code with
 - **make qemu-nox** 

Compiles and run xv6 with qemu
 - qemu-nox run the console version of the emulator

CS 1550 – xv6

```
linux.cs.pitt.edu - PuTTY
(8) thompson $ make qemu-nox
qemu-system-i386 -nographic -drive file=fs.img,index=1,media=disk,format=raw -drive file=xv6.img,index=0,media=disk,format=raw -smp 2 -m 512

(process:118651): GLib-WARNING **: gmem.c:483: custom memory allocation vtable not supported
xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$
```


CS 1550 – xv6

- Once in xv6 you can run **ls**

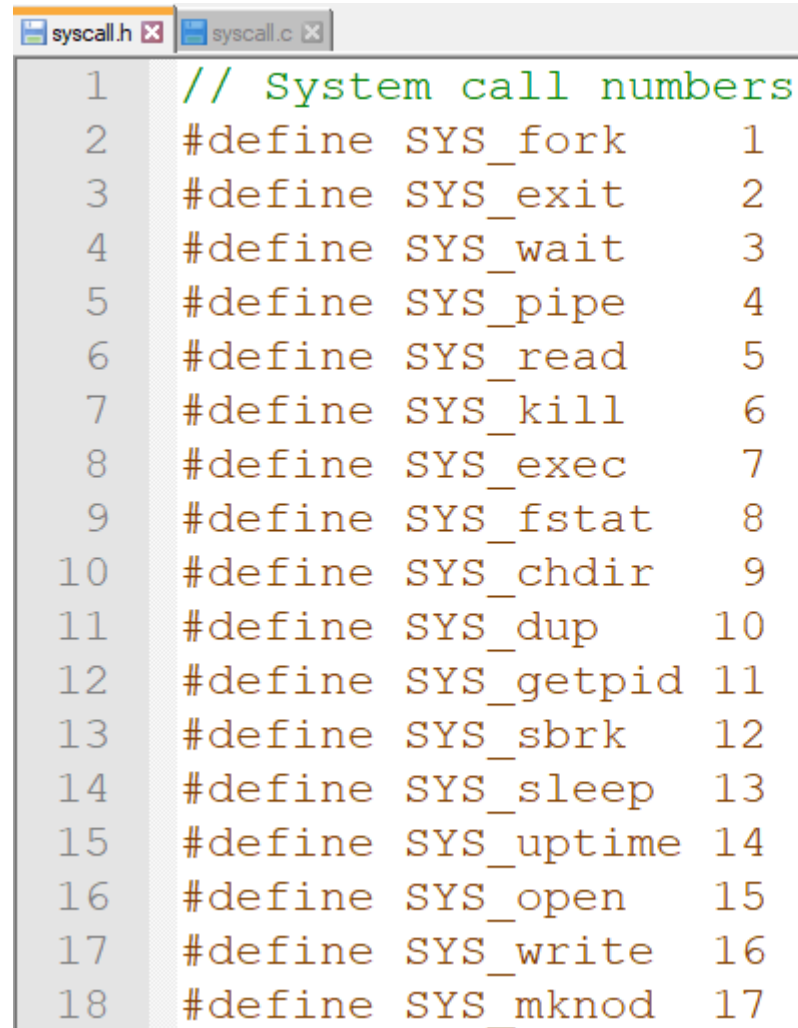
```
cat      2 3 14484
echo     2 4 13340
forktest 2 5 8164
grep     2 6 16020
init     2 7 14232
kill     2 8 13372
ln       2 9 13312
ls       2 10 16172
mkdir    2 11 13404
rm       2 12 13380
sh       2 13 24820
stressfs 2 14 14328
usertests 2 15 67260
wc       2 16 15148
zombie   2 17 13040
console  3 18 0
temp     1 19 32
$
```

CS 1550 – xv6 – Adding a custom Syscall

- In Lab1, you need to add a syscall “getcount”.
- As an example, let’s try to add a syscall named “getday”, which returns a hardcoded pseudo date

CS 1550 – xv6 – Adding a custom Syscall

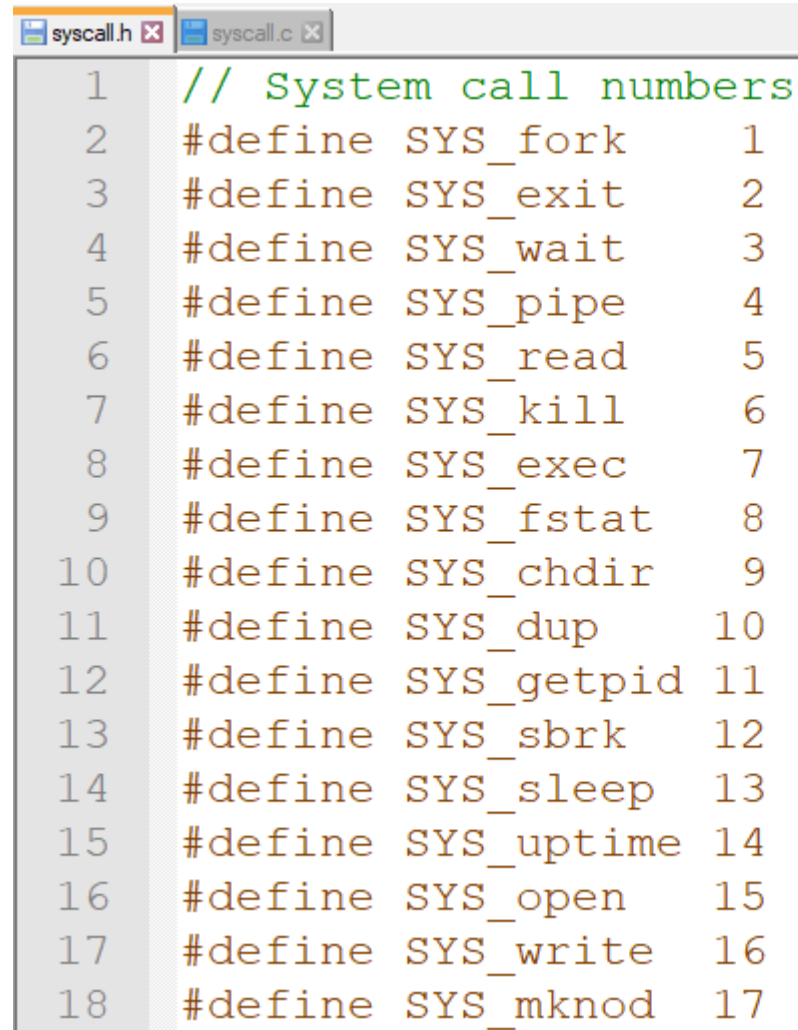
- First, we need to define our new call and its number at
 - **syscall.h**



```
1 // System call numbers
2 #define SYS_fork    1
3 #define SYS_exit    2
4 #define SYS_wait    3
5 #define SYS_pipe    4
6 #define SYS_read    5
7 #define SYS_kill    6
8 #define SYS_exec    7
9 #define SYS_fstat   8
10 #define SYS_chdir   9
11 #define SYS_dup     10
12 #define SYS_getpid  11
13 #define SYS_sbrk    12
14 #define SYS_sleep   13
15 #define SYS_uptime  14
16 #define SYS_open    15
17 #define SYS_write   16
18 #define SYS_mknod   17
```

CS 1550 – xv6 – Adding a custom Syscall

- First, we need to define our new call and its number at
 - `syscall.h`
- Add
 - `#define SYS_getday 22`



The screenshot shows a code editor with two tabs: 'syscall.h' and 'syscall.c'. The 'syscall.h' tab is active, displaying a list of system call definitions. The code is as follows:

```
1 // System call numbers
2 #define SYS_fork    1
3 #define SYS_exit    2
4 #define SYS_wait    3
5 #define SYS_pipe    4
6 #define SYS_read    5
7 #define SYS_kill    6
8 #define SYS_exec    7
9 #define SYS_fstat    8
10 #define SYS_chdir   9
11 #define SYS_dup    10
12 #define SYS_getpid  11
13 #define SYS_sbrk   12
14 #define SYS_sleep  13
15 #define SYS_uptime 14
16 #define SYS_open   15
17 #define SYS_write  16
18 #define SYS_mknod  17
```

CS 1550 – xv6 – Adding a custom Syscall

- Next, we need to map the new call in the array pointer of system calls

- **syscall.c**

- Add

- [SYS_getday] sys_getday,

```
110
111 static int (*syscalls[]) (void) = {
112     [SYS_fork]      sys_fork,
113     [SYS_exit]      sys_exit,
114     [SYS_wait]      sys_wait,
115     [SYS_pipe]      sys_pipe,
116     [SYS_read]      sys_read,
117     [SYS_kill]      sys_kill,
118     [SYS_exec]      sys_exec,
119     [SYS_fstat]     sys_fstat,
120     [SYS_chdir]     sys_chdir,
121     [SYS_dup]       sys_dup,
122     [SYS_getpid]    sys_getpid,
123     [SYS_sbrk]      sys_sbrk,
124     [SYS_sleep]     sys_sleep,
125     [SYS_uptime]    sys_uptime,
126     [SYS_open]      sys_open,
127     [SYS_write]     sys_write,
```

CS 1550 – xv6 – Adding a custom Syscall

- Next, we need to map the new call in the array pointer of system calls
 - **syscall.c**
- Add
 - [SYS_getday] sys_getday,
- Add
 - extern int sys_getday(void);

```
syscall.h x  syscall.c x
94  extern int sys_link(void);
95  extern int sys_mkdir(void);
96  extern int sys_mknod(void);
97  extern int sys_open(void);
98  extern int sys_pipe(void);
99  extern int sys_read(void);
100 extern int sys_sbrk(void);
101 extern int sys_sleep(void);
102 extern int sys_unlink(void);
103 extern int sys_wait(void);
104 extern int sys_write(void);
105 extern int sys_uptime(void);
106
107 static int (*syscalls[])(void) = {
108     [SYS_fork]    sys_fork,
109     [SYS_exit]    sys_exit,
110     [SYS_wait]    sys_wait,
```

CS 1550 – xv6 – Adding a custom Syscall

- Then we need to implement the actual method
- In xv6 this is organized in two files.
 - `sysfile.c` -> file related system calls
 - **`sysproc.c` -> all the other syscalls**

CS 1550 – xv6 – Adding a custom Syscall

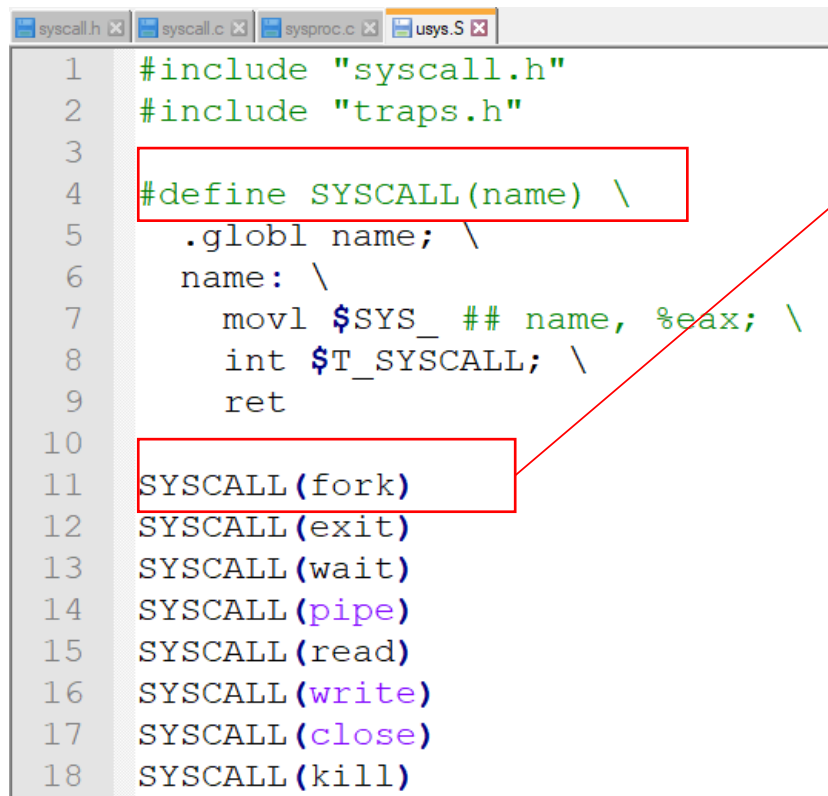
- Then we need to implement the actual method
- In xv6 this is organized in two files.
 - sysfile.c -> file related system calls
 - **sysproc.c -> all the other syscalls**

```
int
sys_getday(void)
{
    return 6;
}
```

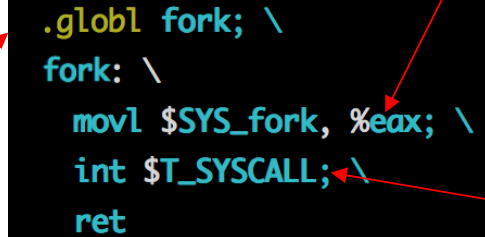
```
syscall.h x syscall.c x sysproc.c x
4  #include "date.h"
5  #include "param.h"
6  #include "memlayout.h"
7  #include "mmu.h"
8  #include "proc.h"
9
10 int
11 sys_fork(void)
12 {
13     return fork();
14 }
15
16 int
17 sys_exit(void)
18 {
19     exit();
20     return 0; // not reached
21 }
```


CS 1550 – xv6 – Adding a custom Syscall

- Afterwards we define the interface for user programs to call
 - Open `usys.S`



```
1 #include "syscall.h"
2 #include "traps.h"
3
4 #define SYSCALL(name) \
5     .globl name; \
6     name: \
7         movl $SYS_ ## name, %eax; \
8         int $T_SYSCALL; \
9         ret
10
11 SYSCALL(fork)
12 SYSCALL(exit)
13 SYSCALL(wait)
14 SYSCALL(pipe)
15 SYSCALL(read)
16 SYSCALL(write)
17 SYSCALL(close)
18 SYSCALL(kill)
```



```
.globl fork; \
fork: \
    movl $SYS_fork, %eax; \
    int $T_SYSCALL; \
    ret
```

Move the syscall ID into `eax` register. The register value will be further pushed to stack (memory).

INT (interrupt) `T_SYSCALL`, which means there's a syscall request

CS 1550 – xv6 – Adding a custom Syscall

- Afterwards we define the interface for user programs to call
 - Open `usys.S`

```
1 #include "syscall.h"
2 #include "traps.h"
3
4 #define SYSCALL(name) \
5     .globl name; \
6     name: \
7         movl $SYS_ ## name, %eax; \
8         int $T_SYSCALL; \
9         ret
10
11 SYSCALL(fork)
12 SYSCALL(exit)
13 SYSCALL(wait)
14 SYSCALL(pipe)
15 SYSCALL(read)
16 SYSCALL(write)
17 SYSCALL(close)
18 SYSCALL(kill)
```

Use this ID to
invoke
corresponding
syscall routine

```
.globl fork; \
fork: \
    movl $SYS_fork, %eax; \
    int $T_SYSCALL; \
    ret
```

Move the syscall ID
into `eax` register. The
register value will be
further pushed to stack
(memory).

INT (interrupt) `T_SYSCALL`,
which means there's a
syscall request

In the end of `syscall.c`

```
void
syscall(void)
{
    int num;
    struct proc *curproc = myproc();

    num = curproc->tf->eax;
    if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
        curproc->tf->eax = syscalls[num]();
    } else {
        cprintf("%d %s: unknown sys call %d\n",
            curproc->pid, curproc->name, num);
        curproc->tf->eax = -1;
    }
}
```

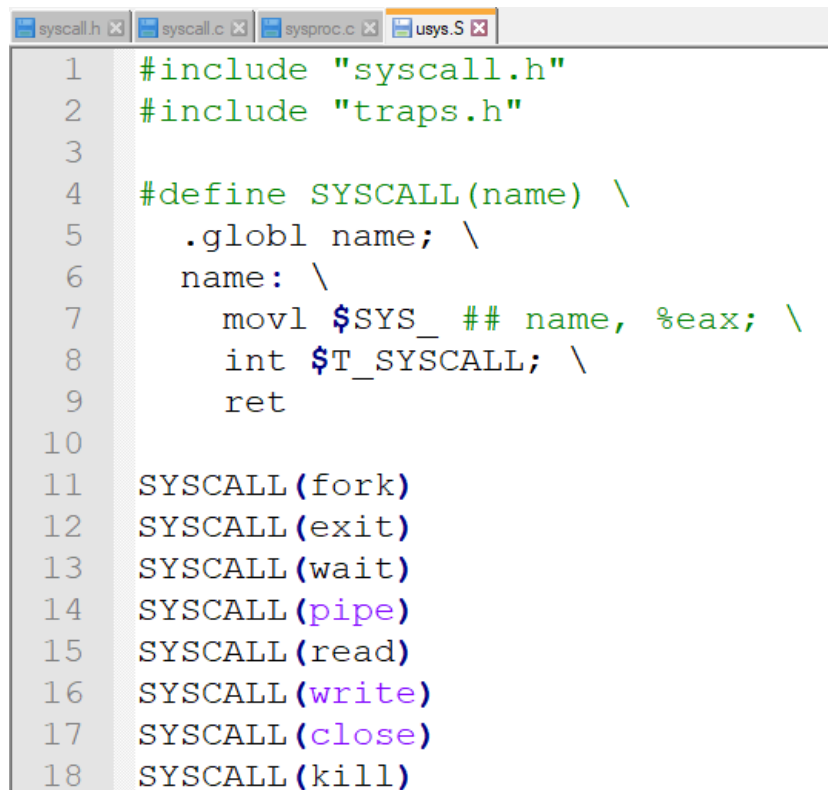
Entrance of all syscall
requests

Each process has its
own per-process
state

Retrieve syscall ID
(now it's already
saved in the stack
related to current
process)

CS 1550 – xv6 – Adding a custom Syscall

- Afterwards we define the interface for user programs to call
 - Open usys.S



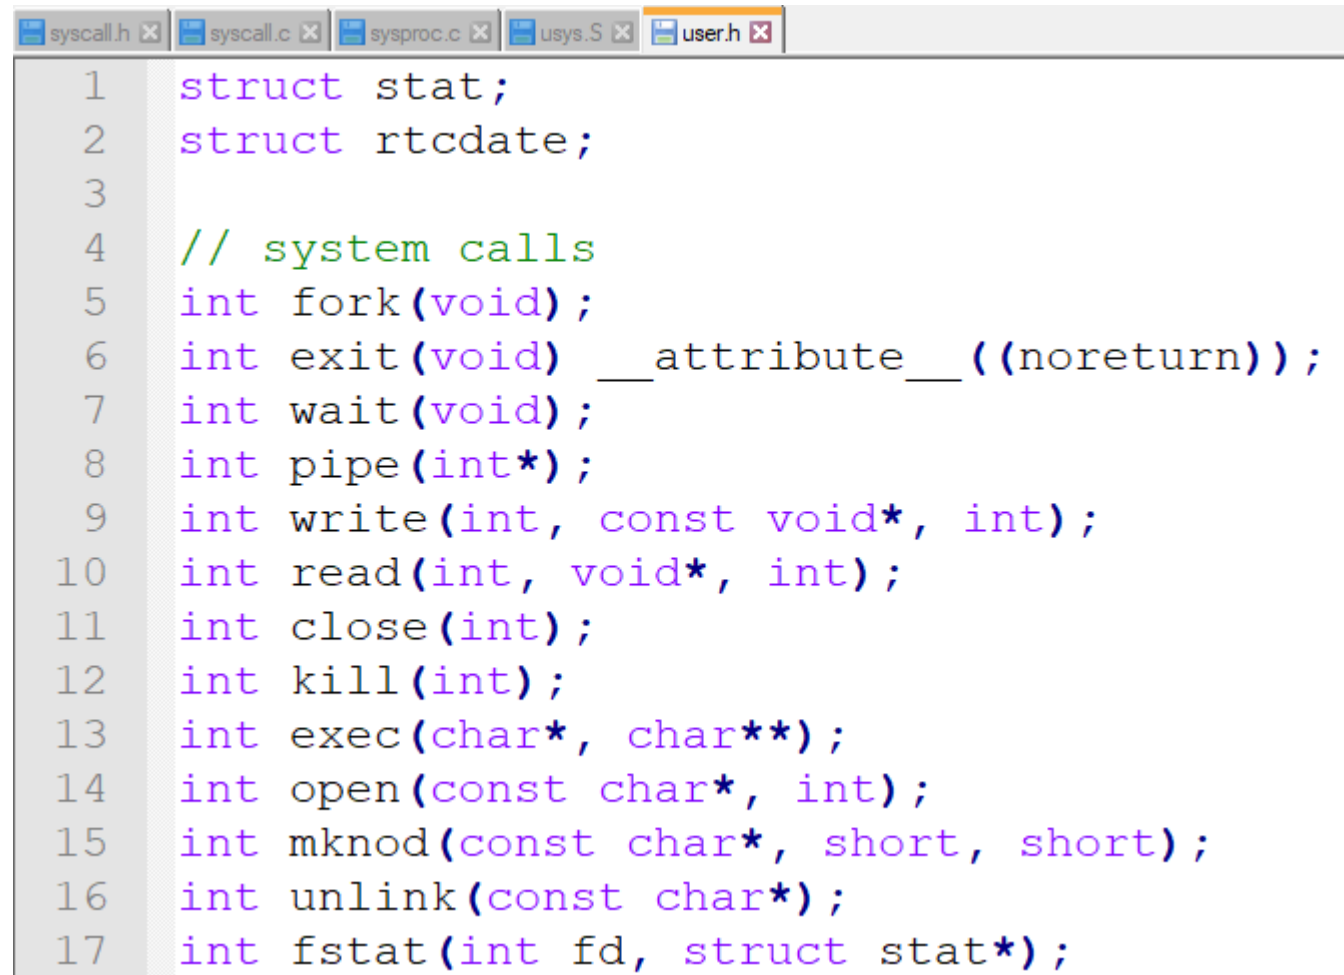
```
1  #include "syscall.h"
2  #include "traps.h"
3
4  #define SYSCALL(name) \
5      .globl name; \
6      name: \
7          movl $SYS_ ## name, %eax; \
8          int $T_SYSCALL; \
9          ret
10
11 SYSCALL(fork)
12 SYSCALL(exit)
13 SYSCALL(wait)
14 SYSCALL(pipe)
15 SYSCALL(read)
16 SYSCALL(write)
17 SYSCALL(close)
18 SYSCALL(kill)
```

What you need to do in usys.S?

Add “SYSCALL(getday)”

CS 1550 – xv6 – Adding a custom Syscall

- Finally we open
 - user.h
- Add “int getday(void);”



```
1 struct stat;
2 struct rtcdate;
3
4 // system calls
5 int fork(void);
6 int exit(void) __attribute__((noreturn));
7 int wait(void);
8 int pipe(int*);
9 int write(int, const void*, int);
10 int read(int, void*, int);
11 int close(int);
12 int kill(int);
13 int exec(char*, char**);
14 int open(const char*, int);
15 int mknod(const char*, short, short);
16 int unlink(const char*);
17 int fstat(int fd, struct stat*);
```

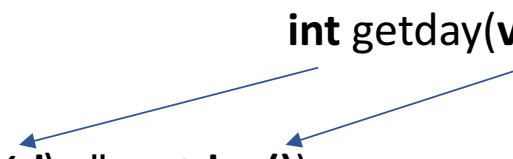
CS 1550 – xv6 – Adding a custom Syscall

- Write a user program for Xv6 to show today's date
 - todays_date.c

```
#include "types.h"
#include "stat.h"
#include "user.h"

int main(void) {
    printf(1, "Today is %d\n", getday());
    exit();
}
```

int getday(void)



The diagram consists of two blue arrows. One arrow originates from the `getday()` argument in the `printf` function call within the `main` function and points to the `int getday(void)` declaration. The second arrow originates from the `getday` part of the `int getday(void)` declaration and points to the `#include "user.h"` line, indicating that the function is defined in that header file.

CS 1550 – xv6 – Adding a custom Syscall

- Adding a user program, such that you can actually run it in XV6
 - Open **Makefile**
- Add
 - `_today's_date\`

```
syscall.h x syscall.c x sysproc.c x usys.S x user.h x today's_date.c x Makefile x
166 .PRECIOUS: %.o
167
168 UPROGS=\
169     _cat\
170     _echo\
171     _forktest\
172     _grep\
173     _init\
174     _kill\
175     _ln\
176     _ls\
177     _mkdir\
178     _rm\
179     _sh\
180     _stressfs\
181     _usertests\
182     _wc\
183     _zombie\
184
```

cat	2	3	14484
echo	2	4	13340
forktest	2	5	8164
grep	2	6	16020
init	2	7	14232
kill	2	8	13372
ln	2	9	13312
ls	2	10	16172
mkdir	2	11	13404
rm	2	12	13380
sh	2	13	24820
stressfs	2	14	14328
usertests	2	15	67260
wc	2	16	15148
zombie	2	17	13040
console	3	18	0
temp	1	19	32
\$			

CS 1550 – xv6 – Adding a custom Syscall

- Adding a user program
 - Open **Makefile**
- and add
 - `today's_date.c\`

```
250
251 EXTRA=\
252     mkfs.c ulib.c user.h cat.c echo.c forktest.c grep.c kill.c\
253     ln.c ls.c mkdir.c rm.c stressfs.c usertests.c wc.c zombie.c\
254     printf.c umalloc.c\
255     README dot-bochsrc *.pl toc.* runoff runoffl runoff.list\
256     .gdbinit.tmpl gdbutil\
257
258
259 dist:
260     rm -rf dist
261     mkdir dist
262     for i in $(FILES); \
263     do \
```

CS 1550 – xv6 – Adding a custom Syscall

- Adding a user program
 - Open makefile
- and add
 - `todays_date.c\`
- Finally execute
 - `make qemu-nox`
- Run the program by calling it directly within xv6
 - `todays_date`

CS 1550 – xv6 exercise hints

- We need to worry about two things:
 - How to count syscalls invoked by the calling process?
 - Implement the method to return the count number given a syscall ID

CS 1550 – xv6 exercise hints

- Syscall calls will need a variable to hold the counting values
 - Where to write this data structure?
 - Which file holds processes metadata (per-process state)? `proc.c`, `proc.h`
 - The key function can be found in `syscall.c`
 - `syscall(void)` -> Is called every time any syscall is called

CS 1550 – xv6 exercise hints

```
void
syscall(void)
{
    int num;
    struct proc *curproc = myproc();

    num = curproc->tf->eax;
    if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
        curproc->tf->eax = syscalls[num]();
    } else {
        cprintf("%d %s: unknown sys call %d\n",
            curproc->pid, curproc->name, num);
        curproc->tf->eax = -1;
    }
}
```

The system call numbers match the entries in the syscalls array, a table of function pointers



CS 1550 – xv6 exercise hints

- Implementing **getcount**
 - Specify the method and its **id** in **syscall.h**
 - Specify extern method and pointer
 - **syscall.c**
 - Where to implement int **sys_getcount(void)**?
 - **sysproc.c**
 - Add SYSCALL(getcount)
 - **usys.S**
 - **getcount.c**
 - Modify **proc.h**, **proc.c** according to your method of counting.
 - Declare counting array? Hints: check **proc.h** to find the per-process state structure
 - Initialize counting array? Hints: check **proc.c** to find the function that performs initialization for processes

CS 1550 – xv6 exercise hints

- Submit to GradeScope the files that you have modified within the source code of xv6.
- You should modify the following files only:
 - syscall.h
 - syscall.c
 - user.h
 - usys.S
 - proc.h
 - proc.c
 - sysproc.c
 - Makefile