



CS 1550


Week 2 – Setting up for Project 1
& Lab 1 cont.

TA: Jinpeng Zhou

jiz150@pitt.edu

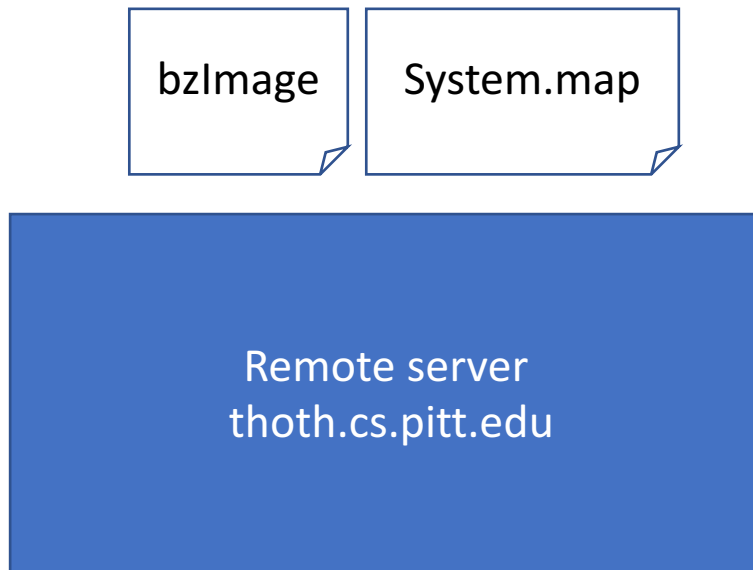
Setting up environment for Project 1

- We will modify **Linux** source code and build a Linux distro on thoth server, then run & test our build on a VM



Setting up environment for Project 1

- We need to build the new kernel on **thoth server (thoth.cs.pitt.edu)**




Setting up environment for Project 1

- We need to run our build on a QEMU Virtual Machine (VM), e.g., a VM on your local machine or on thoth
 - 1st, launch the default VM
 - 2nd, download our build (bzImage, System.map) to the VM
 - 3rd, configure the VM such that we can reboot it to our own kernel

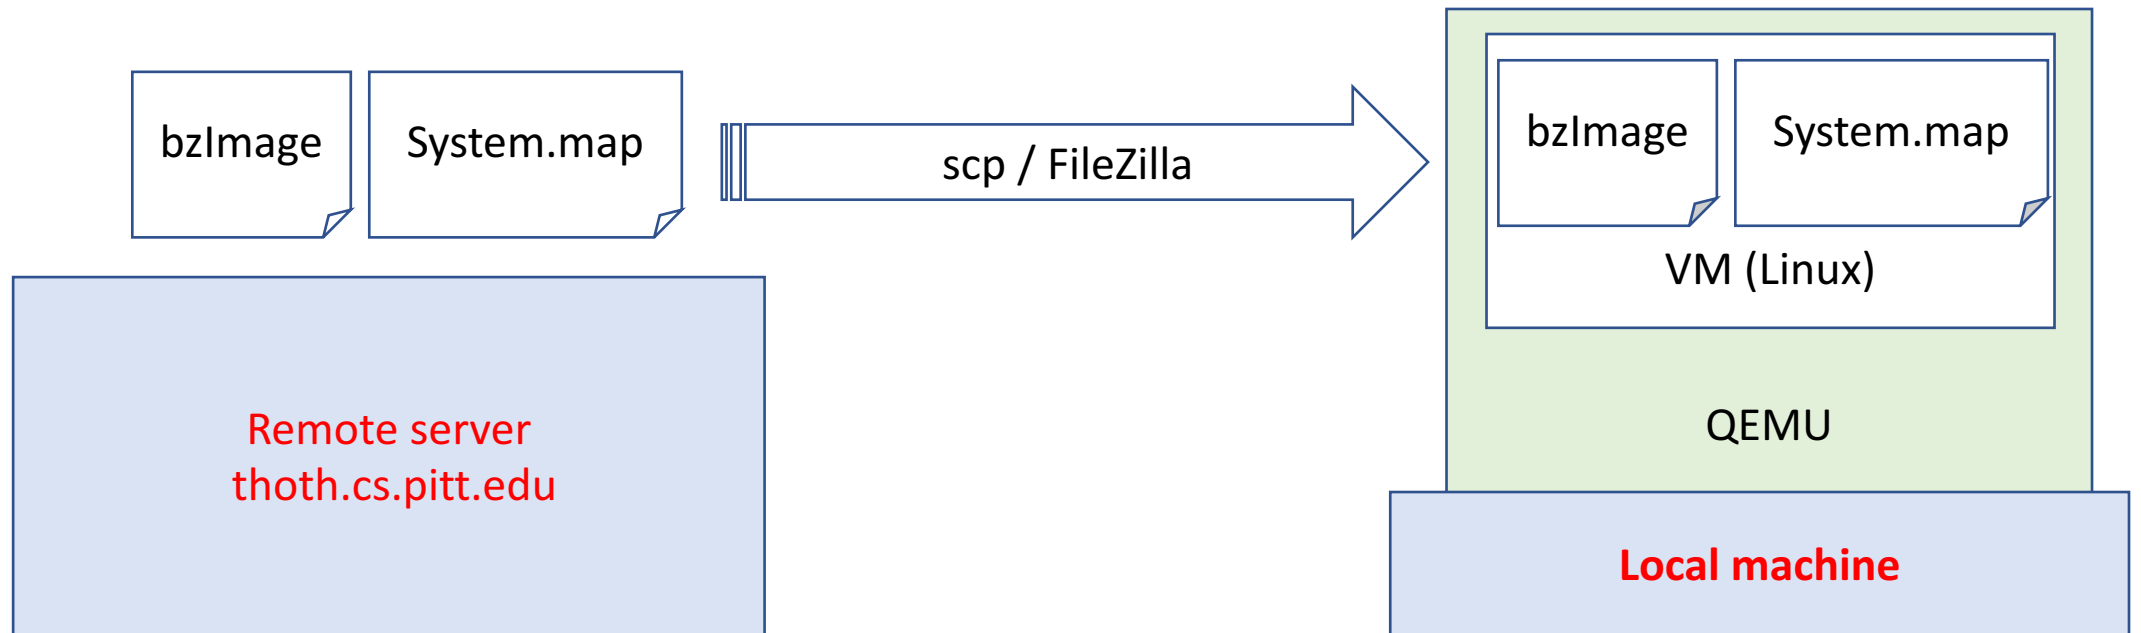
Setting up environment for Project 1

- 1st, **launch** the default VM: install qemu, then launch the VM image.




Setting up environment for Project 1

- 2nd, **download** our build (bzImage, System.map) to the VM



Setting up environment for Project 1

- 3rd, **configure** the VM such that we can **reboot** it with our build (a modified linux)



Setting up environment for Project 1

1. First log to your **thoth.cs.pitt.edu** account
 - SSH

Setting up environment for Project 1

1. First log to your **thoth.cs.pitt.edu** account
 - SSH
2. Navigate to **/u/OSLab/username**
 - “/u/OSLab/username” will be your working directory. You can also use “~/private” as your working directory.
 - “cd /u/OSLab/ABC123” or “cd ~/private”

```
THIS SYSTEM IS FOR THE USE OF AUTHORIZED USERS ONLY.

Individuals using this computer system without authority, or in
excess of their authority, are subject to having all of their
activities on this system monitored and recorded by system
personnel.

In the course of monitoring individuals improperly using this
system, or in the course of system maintenance, the activities
of authorized users may also be monitored.

Anyone using this system expressly consents to such monitoring
and is advised that if such monitoring reveals possible
evidence of criminal activity, system personnel may provide the
evidence of such monitoring to law enforcement officials.

(1) thoth $ ls
Backup bin News perl5 private public
(2) thoth $
```

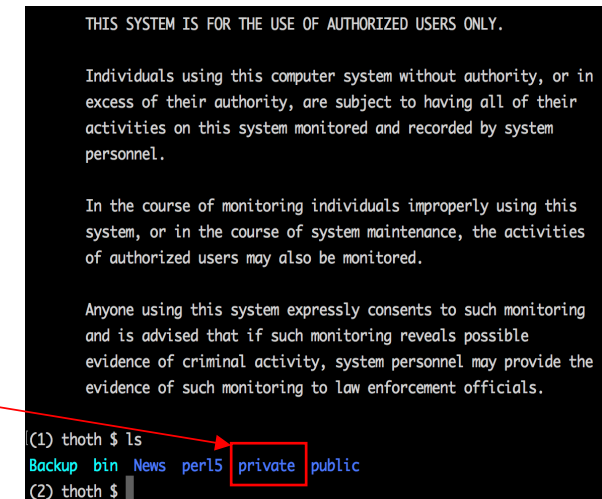
Setting up environment for Project 1

1. First log to your **thoth.cs.pitt.edu** account

- SSH

2. Navigate to **/u/OSLab/username**

- “/u/OSLab/username” will be your working directory. You can also use “~/private” as your working directory.
 - “cd /u/OSLab/ABC123” or “cd ~/private”
- Copy linux source from /u/OSLab/original/linux-2.6.23.1.tar.bz2
 - **cp /u/OSLab/original/linux-2.6.23.1.tar.bz2 .**



THIS SYSTEM IS FOR THE USE OF AUTHORIZED USERS ONLY.

Individuals using this computer system without authority, or in excess of their authority, are subject to having all of their activities on this system monitored and recorded by system personnel.

In the course of monitoring individuals improperly using this system, or in the course of system maintenance, the activities of authorized users may also be monitored.

Anyone using this system expressly consents to such monitoring and is advised that if such monitoring reveals possible evidence of criminal activity, system personnel may provide the evidence of such monitoring to law enforcement officials.

```
(1) thoth $ ls
Backup bin News perl5 private public
(2) thoth $
```

A red arrow points from the text “~/private” in the instructions to the “private” directory in the terminal output.

This dot means current position. Since we just run “cd”, it will be inside the working directory

Setting up environment for Project 1

3. Extract files locally

- Run `tar xjf linux-2.6.23.1.tar.bz2`

Setting up environment for Project 1

3. Extract files locally

- Run `tar xvj linux-2.6.23.1.tar.bz2`

4. Enter into **linux-2.6.23.1/**

- Run `cd linux-2.6.23.1`

Setting up environment for Project 1

3. Extract files locally

- Run `tar xjf linux-2.6.23.1.tar.bz2`


4. Enter into **linux-2.6.23.1/**

- Run `cd linux-2.6.23.1`

5. Copy the .config file

- Run `cp /u/OSLab/original/.config .`

Current position: WORKING_DIR/linux-2.6.23.1



Setting up environment for Project 1

3. Extract files locally

- Run `tar xjf linux-2.6.23.1.tar.bz2`

4. Enter into **linux-2.6.23.1/**

- Run `cd linux-2.6.23.1`

5. Copy the .config file

- Run `cp /u/OSLab/original/.config .`

6. Build linux source code

- Run `make ARCH=i386 bzImage`

7. Compile test programs

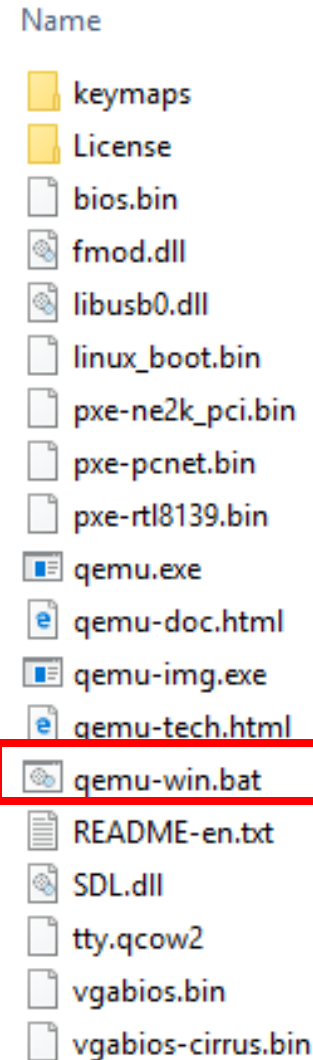
- Run `gcc -m32 -o trafficsim -I /u/OSLab/USERNAME/linux-2.6.23.1/include/ trafficsim.c`

Setting up environment for Project 1

- After change the kernel code, we need to rebuild the Kernel
 - Run again `make ARCH=i386 bzImage`

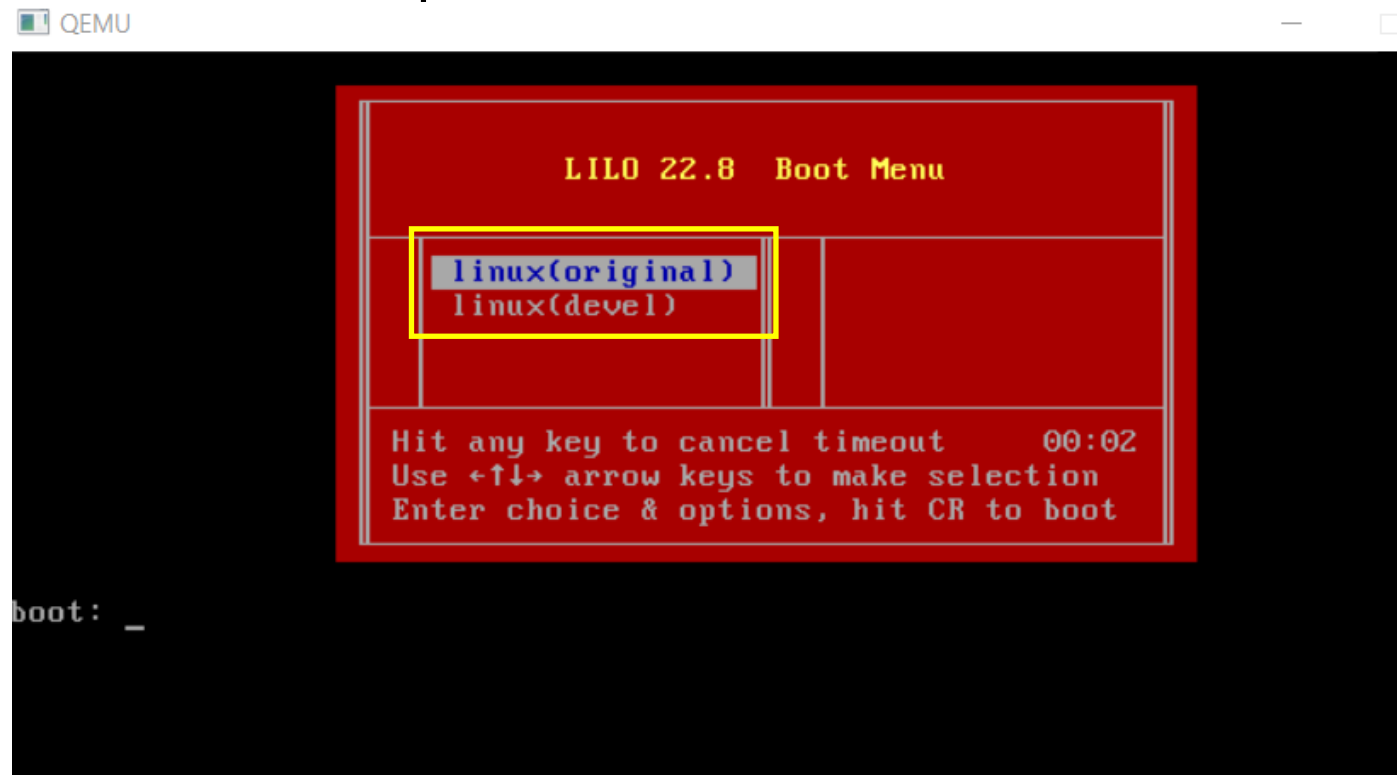
Configuring QEMU

- We need a x86 version of QEMU (username and pass is **root**). It will be available in the “Assignment” on canvas.
- Windows:
 - double-click/execute qemu-win.bat
- Mac/Ubuntu
 - Run the script:
./start.sh



Configuring QEMU

- Choose Linux(original): we will boot into “original”, download our build, connect it to the “devel”, then reboot to “devel”
 - User ‘**root**’ as user and password



Copying files from Linux to QEMU

- Now we need two files from the Linux we just built
 - Kernel File **bzImage** from:
 - linux- 2.6.23.1/arch/i386/boot/
 - System call map **System.map** from:
 - linux-2.6.23.1/
- Please be sure about **the path** where **you copied** the **linux distro!**
 - If you follow the steps here the linux files should be in **/u/OSLab/username**

Copying files from Linux to QEMU

- **FROM WITHIN THE NEW QEMU**

- Now, we are in the default directory (in vm) after we type in “root” as user&pwd

- Download the files from your compiled Linux:

- **scp** USERNAME@thoth.cs.pitt.edu:/u/OSLab/username/linux-2.6.23.1/arch/i386/boot/**bzImage** .
 - **scp** USERNAME@thoth.cs.pitt.edu:/u/OSLab/username/linux-2.6.23.1/**System.map** .

Dot (current position):
The default directory

- Install the rebuilt kernel in QEMU (to connect our build to “devel”):

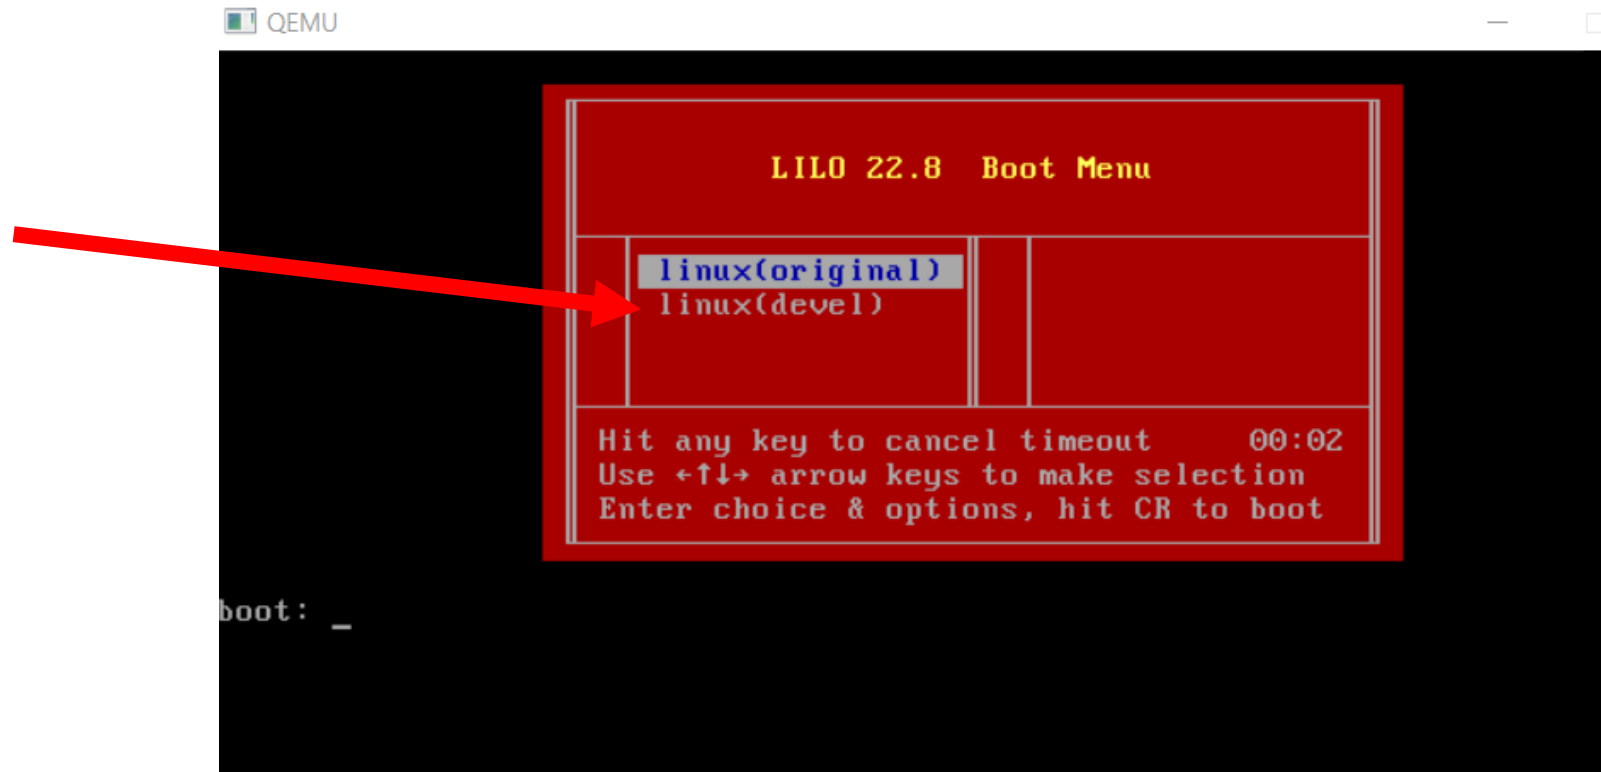
- cp bzImage /boot/bzImage-devel
 - cp System.map /boot/System.map-devel

Copying files from Linux to QEMU

- After this run **linux loader** command:
 - Run **lilo**
 - This will relink the new modified kernel you just copied
- Then reboot the system with the command:
 - Run **reboot**

Copying files from Linux to QEMU

- You will change to **linux(devel)** kernel
 - So to see changes always remind to choose it when opening Qemu



Run test program

Follow the same process:

- 1st, build the test program **trafficsim on thoth server**
 - `gcc -g -m32 -o trafficsim -I/u/OSLab/USERNAME/linux-2.6.23.1/include/ trafficsim.c`
 - `gcc` `-g` `-m32` `-o` `trafficsim` `-I/u/OSLab/USERNAME/linux-2.6.23.1/include/` `trafficsim.c`
 - `Exec name` `Header-file path` `Src file`
 - For different tests, you only need to change src file and target executable name.
- 2nd, download the test program **to VM (“devel”)**
 - After boot to the modified linux (“devel”), run scp to download tests
- 3rd, run it **inside VM(“devel”)**
 - `./trafficsim`

Use GDB for project

You can run qemu on thoth server, then debug your kernel on thoth.

- Set up qemu on thoth
 1. `cd /u/OSLab/USERNAME`
 2. `cp /u/OSLab/original/qemu.tar.gz .`
 3. `tar xzvf qemu.tar.gz`
 4. `cd qemu`
 5. `make qemu`

Use GDB for project

You can run qemu on thoth server, then debug your kernel on thoth.

- We'll need two connections (two terminals)
 - One is to launch the VM
 - The other is to launch gdb to debug the running VM

Use GDB for project

You can run qemu on thoth server, then debug your kernel on thoth.

- Launch the VM in first terminal
 - 1st, copy the modified kernel from your working directory to qemu folder
 - `cd /u/OSLab/USERNAME/qemu`
 - `cp /u/OSLab/USERNAME/linux-2.6.23.1/vmlinux .` ← Current position:
/u/OSLab/USERNAME/qemu
 - 2nd, launch vm with gdb enabled
 - `make qemu-gdb`
 - You'll see the vm is halted by gdb upon boot:



VGA Blank mode

Use GDB for project

You can run qemu on thoth server, then debug your kernel on thoth.

- Launch the gdb in second terminal
 - 1st, open a new SSH connection to thoth.
 - 2nd, launch gdb
 - `cd /u/OSLab/USERNAME/qemu`
 - `gdb -iex "set auto-load safe-path ."`

```
(5) thoth $ gdb -iex "set auto-load safe-path ."  
GNU gdb (GDB) Red Hat Enterprise Linux (7.2-64.el6_5.2)  
Copyright (C) 2010 Free Software Foundation, Inc.  
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl>  
This is free software: you are free to change and redistribute it.  
There is NO WARRANTY, to the extent permitted by law. Type "show copy  
and "show warranty" for details.  
This GDB was configured as "x86_64-redhat-linux-gnu".  
For bug reporting instructions, please see:  
<http://www.gnu.org/software/gdb/bugs/>.  
The target architecture is assumed to be i386  
+ target remote localhost:29209  
0x0000ffff in ?? ()  
+ symbol-file vmlinux
```

Use GDB for project

You can run qemu on thoth server, then debug your kernel on thoth.

- Launch the gdb in second terminal
 - 1st, open a new SSH connection to thoth.
 - 2nd, launch gdb
 - `cd /u/OSLab/USERNAME/qemu`
 - `gdb -iex "set auto-load safe-path ."`
- You'll see gdb has attached to the vm
 - Now, you can use gdb operations, e.g.:
 - Set a breakpoint by `"b sys_cs1550_down"`
 - Continue by `"c"`
 - The halted vm (in first terminal) will continue and finish booting
 - If you run the traffisim inside the VM. It will trigger the invocation of `sys_cs1550_down`, then the kernel will be halted upon `sys_cs1550_down`

```
(5) thoth $ gdb -iex "set auto-load safe-path ."  
GNU gdb (GDB) Red Hat Enterprise Linux (7.2-64.el6_5.2)  
Copyright (C) 2010 Free Software Foundation, Inc.  
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl  
This is free software: you are free to change and redistribute it.  
There is NO WARRANTY, to the extent permitted by law. Type "show copy  
and "show warranty" for details.  
This GDB was configured as "x86_64-redhat-linux-gnu".  
For bug reporting instructions, please see:  
<http://www.gnu.org/software/gdb/bugs/>.  
The target architecture is assumed to be i386  
+ target remote localhost:29209  
0x0000fff0 in ?? ()  
+ symbol-file vmlinux  
(gdb) b sys_cs1550_down  
Breakpoint 1 at 0xc011dd92  
(gdb) c  
Continuing.
```

Use GDB for project

Debugging the user program trafficsim.

- 1st, we need to load the symbol table of trafficsim to gdb
 - Get the text section address of trafficsim (by running readelf in thoth terminal)
 - `readelf -WS PATH_TO_TEST_EXEC/trafficsim | grep text`
 - You'll see an output like this:

```
[13] .text          PROGBITS          08048610 000610 000e1c 00  AX  0  0 16
```

This is the address we need

Use GDB for project

Debugging the user program trafficsim.

- 1st, we need to load the symbol table of trafficsim to gdb
 - Get the text section address of trafficsim (by running readelf in thoth terminal)

- `readelf -WS PATH_TO_TEST_EXEC/trafficsim | grep text`

- You'll see an output like this:

```
[13] .text          PROGBITS          08048610 000610 000e1c 00  AX  0  0 16
```

This is the address we need

- Load the symbol table to gdb (by running add-symbol-table inside gdb)

- (gdb) `add-symbol-file PATH_TO_TEST_EXEC/trafficsim 0x08048610`

- Now, you'll be able to add breakpoint to traffisim.c, e.g.:

- (gdb) `b trafficsim.c:main`

- After that, if you run the trafficsim inside the vm, it will be halted by gdb at its main()



Lab 1 cont.

Add a new syscall to XV6

CS 1550 – xv6 – Adding a custom Syscall

Implement syscall inside kernel space

New syscall number (syscall.h): `#define SYS_getcount 22`

Extend syscall table (syscall.c): `[SYS_getday] sys_getcount,`
`extern int sys_getcount(void);`

Implement syscall routine (sysproc.c): `int sys_getcount(void) {`
`// return how many times`
`// this syscall has been called`
`// by the calling process`
`}`

Provide user interface for user-space

User interface definition (usys.S): `SYSCALL(getcount)`

User interface declaration (user.h): `int getcount(int);`

CS 1550 – xv6 – Adding a custom Syscall

Implement syscall inside kernel space

New syscall number (syscall.h): `#define SYS_getcount 22`

Extend syscall table (syscall.c): `[SYS_getday] sys_getcount,`
`extern int sys_getcount(void);`

Implement syscall routine (sysproc.c): `int sys_getcount(void) {`
`// return how many times`
`// this syscall has been called`
`// by the calling process`

`/* which syscall? which process? */`
`}`

CS 1550 – xv6 – Adding a custom Syscall

Implement syscall inside kernel space


- New syscall number (syscall.h): `#define SYS_getcount 22`
- Extend syscall table (syscall.c):
 - `[SYS_getday] sys_getcount,`
 - `extern int sys_getcount(void);`
- Implement syscall routine (sysproc.c): `int sys_getcount(void) {`

User process will call getcount with one parameter: the syscall number N, to get how many times this syscall-N has been called by this user process.

But, inside kernel, all syscall routines takes `void` as parameter, so we need to retrieve the actual parameter (syscall number), by using `argint()`

```
    // return how many times
    // this syscall has been called
    // by the calling process
    /* 1st, identify the calling process */
    struct proc * calling_proc = myproc();
    /* 2nd, retrieve the syscall num */
    int num;
    argint(0, &num);
    /* 3rd, return the corresponding counter */
    .....
}
```

CS 1550 – xv6 – Adding a custom Syscall




Each process will have its own counters, recording the calling times for each syscall. E.g., xv6 by default has 21 syscalls, so each process needs 21 counters.

These counters should be per-process. So we define it as part of per-process metadata: **“struct proc” inside “proc.h”**. After adding this new metadata, we also need to initialize it. We can initialize it upon the process creation: **check the “allocproc()” function inside “proc.c”, and find the position to initialize it**

After that, we are able to **update** and **return** these counters during runtime

```
    // return how many times
    // this syscall has been called
    // by the calling process
    /* 1st, identify the calling process */
    struct proc * calling_proc = myproc();
    /* 2nd, retrieve the syscall num */
    int num;
    argint(0, &num);
    /* 3rd, return the corresponding counter */
    .....
}
```

CS 1550 – xv6 – Adding a custom Syscall



After that, we are able to update and return these counters during runtime

Where can we update the counters??

--Check the logic of the global syscall entrance function:
“syscall()” inside “syscall.c”

This is the position we actually call a syscall routine

```
void
syscall(void)
{
    int num;
    struct proc *curproc = myproc();

    num = curproc->tf->eax;
    if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
        curproc->tf->eax = syscalls[num]();
    } else {
        cprintf("%d %s: unknown sys call %d\n",
            curproc->pid, curproc->name, num);
        curproc->tf->eax = -1;
    }
}
```

CS 1550 – xv6 – Adding a custom Syscall

Test program (getcount.c) is provided in lab1 pdf:

```
main() {  
    ...  
    getcount(N)  
    ...  
}
```

1. Integrate it into XV6 by modifying makefile: UPROGS, EXTRA
2. Rebuild and launch xv6, run “ls”, now you can see the test program listed as a command. Run it and check the results. Expected results are given in the lab1 pdf.

****:** If you saw correct numbers from your own test but wrong numbers from autograder’s test. Most likely, the initialization of per-process counters is wrong. **Such inconsistency typically results from missing or improper initializations. If you met this problem in future labs/projects, check initializations first.**

****:** You can also use gdb to debug your labs. Similar as debugging projects, you need two connections. Please refer to gdb instructions in lab1 pdf.